

Transition to Transport (TCP)

CSE 118: Computer Networks

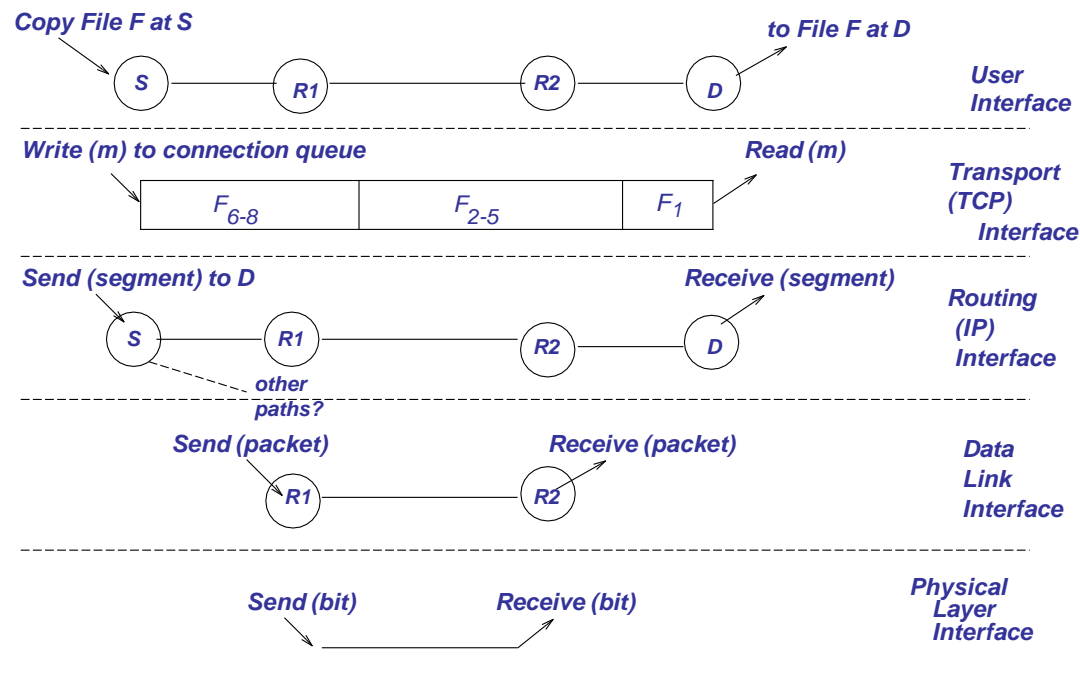
George Varghese

(based on slides by Alex Snoeren)




THE IP ABSTRACTIONS:

Each layer provides a service to the layer above it





Recall Big Picture

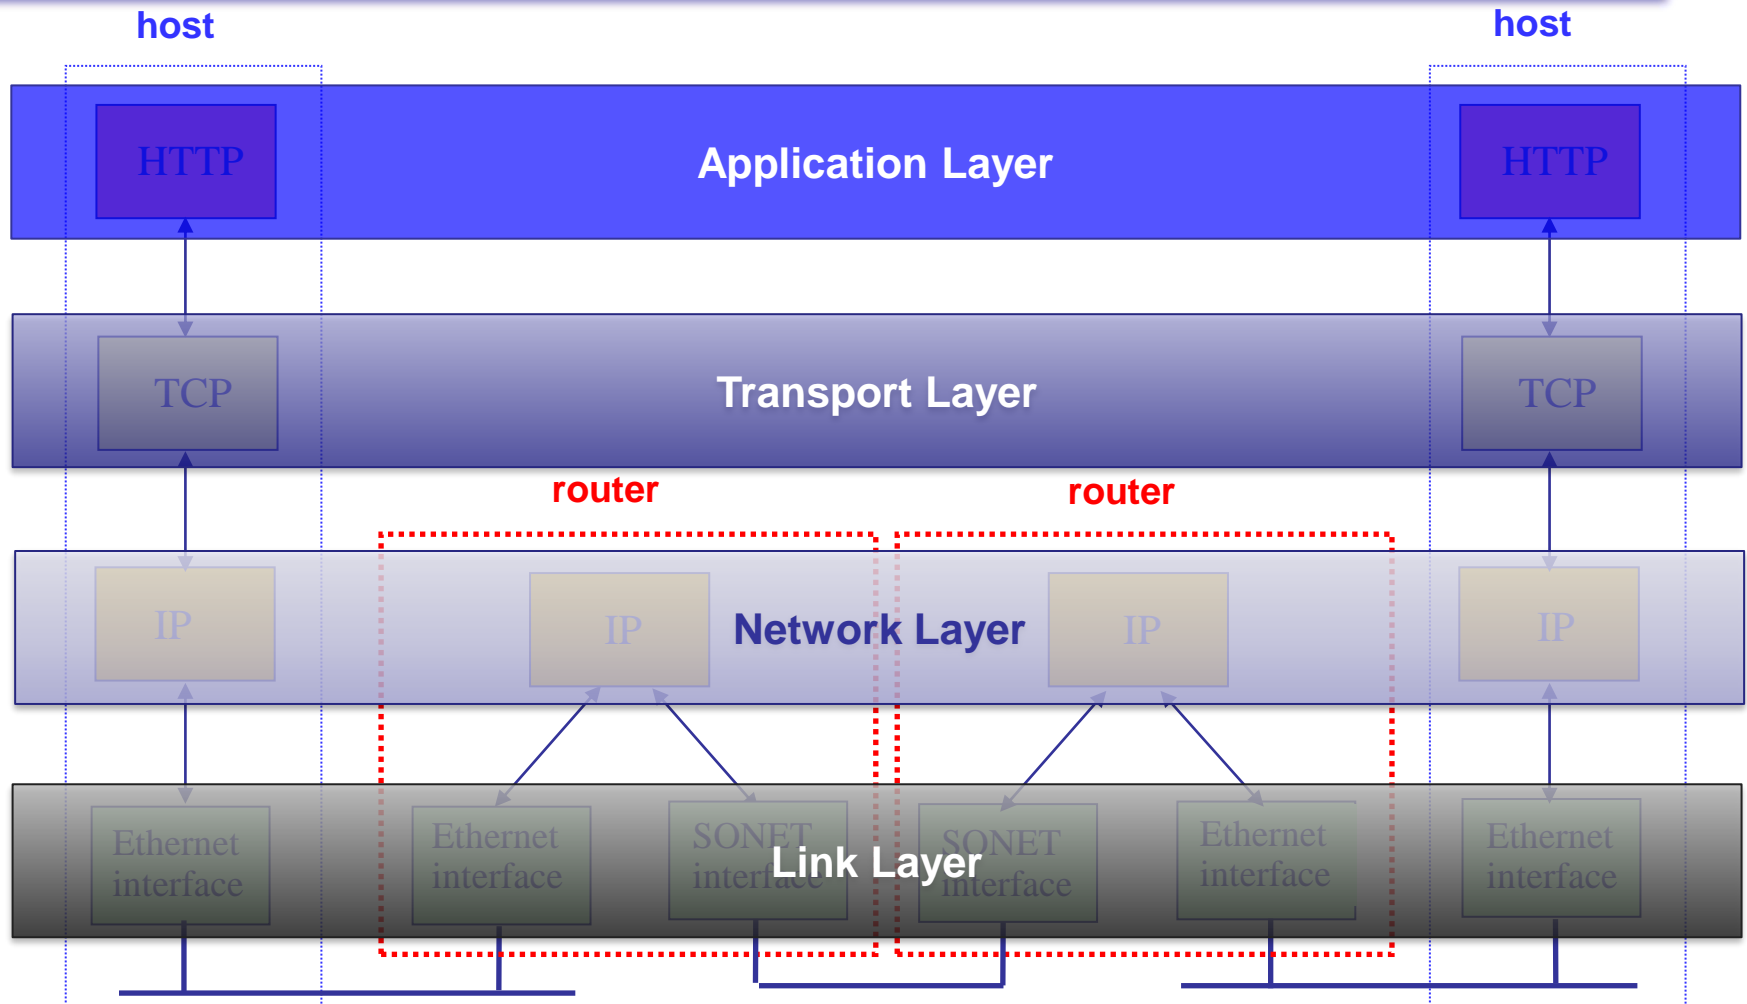


Overview



- Process naming/demultiplexing
- User Datagram Protocol (UDP)
- Transport Control Protocol (TCP)
 - ◆ Three-way handshake
 - ◆ Flow control

Where Transport Layer Fits




Basic Transport Questions



BASIC QUESTIONS

- 1) What function does a transport provide?
- 2) What is a connection?
- 3) Why not have just one connection for all data?
- 4) Why not keep connections up always?
- 5) How do we address the receiving process in the receiver machine?




P1,...P4 denote processes (mail, ftp)
--- lines denote connections

Basic Transport Answers



BASIC QUESTIONS

- 1) Reliable or Unreliable data Delivery between processes
- 2) Shared state for each process pair that enables delivery
- 3) Fast processes would be hostage to slow processes
- 4) Too many possible process pairs \rightarrow close + set up connections
- 5) Need an OS independent mechanism: ports



P1,...P4 denote processes (mail, ftp)
--- lines denote connections



Transport Layer Tasks

- Multiplexing (UDP does only this, so does TCP)
- Reliability (TCP only)
- Flow Control (TCP only)
- Congestion Control (TCP only)

Naming Processes/Services



- Process here is an abstract term for your Web browser (HTTP), Email servers (SMTP), hostname translation (DNS)
- How do we identify for remote communication?
 - ◆ Process id or memory address are OS-specific and transient
- So TCP and UDP use **ports**
 - ◆ 16-bit integers representing mailboxes that processes “rent”
 - ◆ Identify process uniquely as (IP address, protocol, port)



Picking Port Numbers

- We still have the problem of allocating port numbers
 - ◆ What port should a Web server use on host *X*?
 - ◆ To what port should you send to contact that Web server?

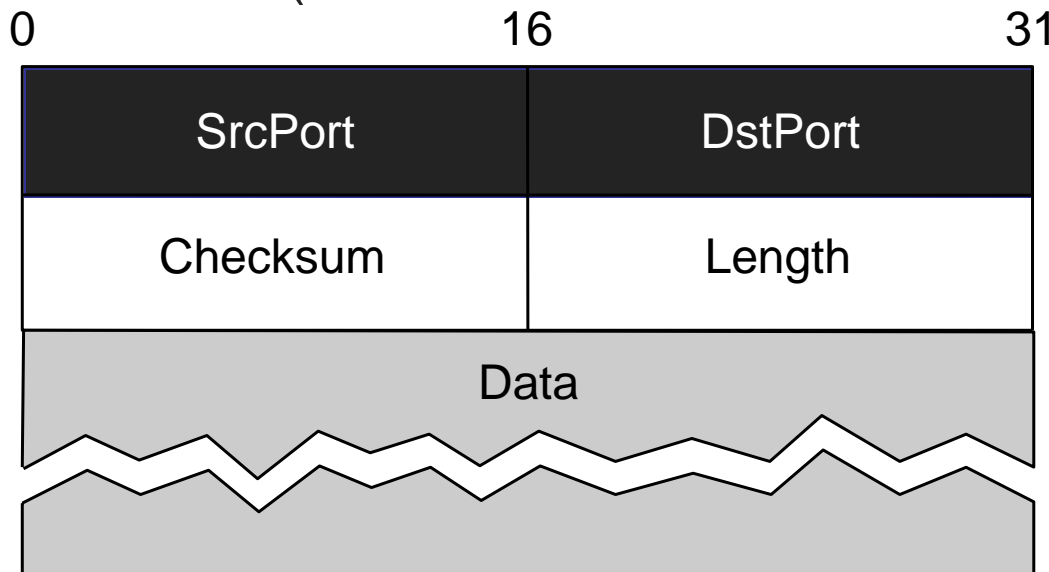
- Servers typically bind to **well-known** port numbers
 - ◆ e.g., HTTP 80, SMTP 25, DNS 53, ... look in /etc/services
 - ◆ Ports below 1024 traditionally reserved for well-known services

- Clients use OS-assigned temporary (**ephemeral**) ports
 - ◆ Above 1024, recycled by OS when client finished

User Datagram Protocol (UDP)




- ▣ Provides *unreliable message delivery* between processes. So what does it do? Multiplexing!
 - ◆ Source port filled in by OS as message is sent
 - ◆ Destination port identifies UDP delivery queue at endpoint
- ▣ Connectionless (no state about who talks to whom)





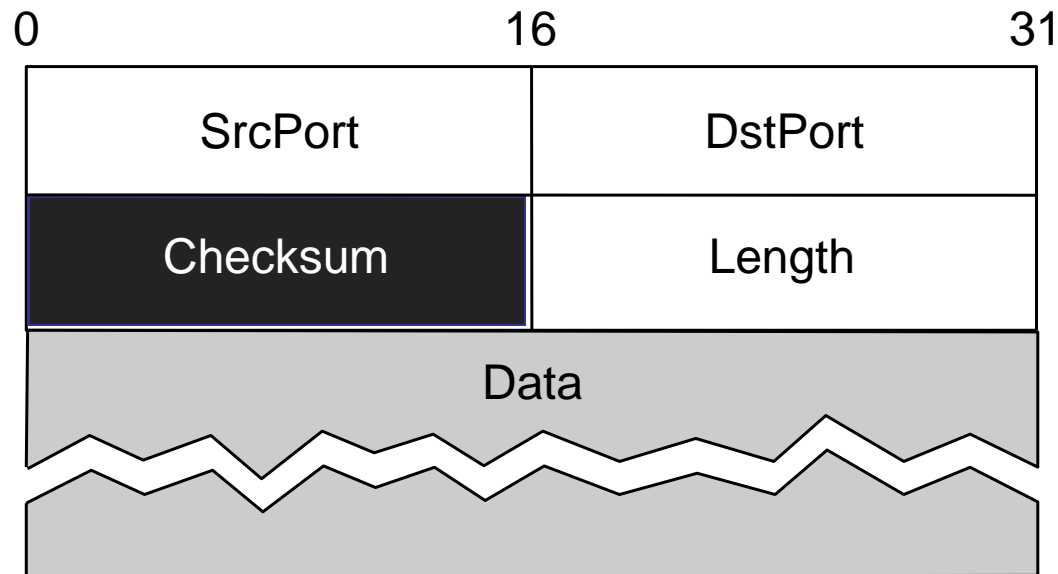
UDP Delivery





UDP Checksum

- UDP includes optional protection against errors
 - ◆ Checksum intended as an end-to-end check on delivery
 - ◆ So it covers data, UDP header, and **IP pseudoheader**





Applications for UDP

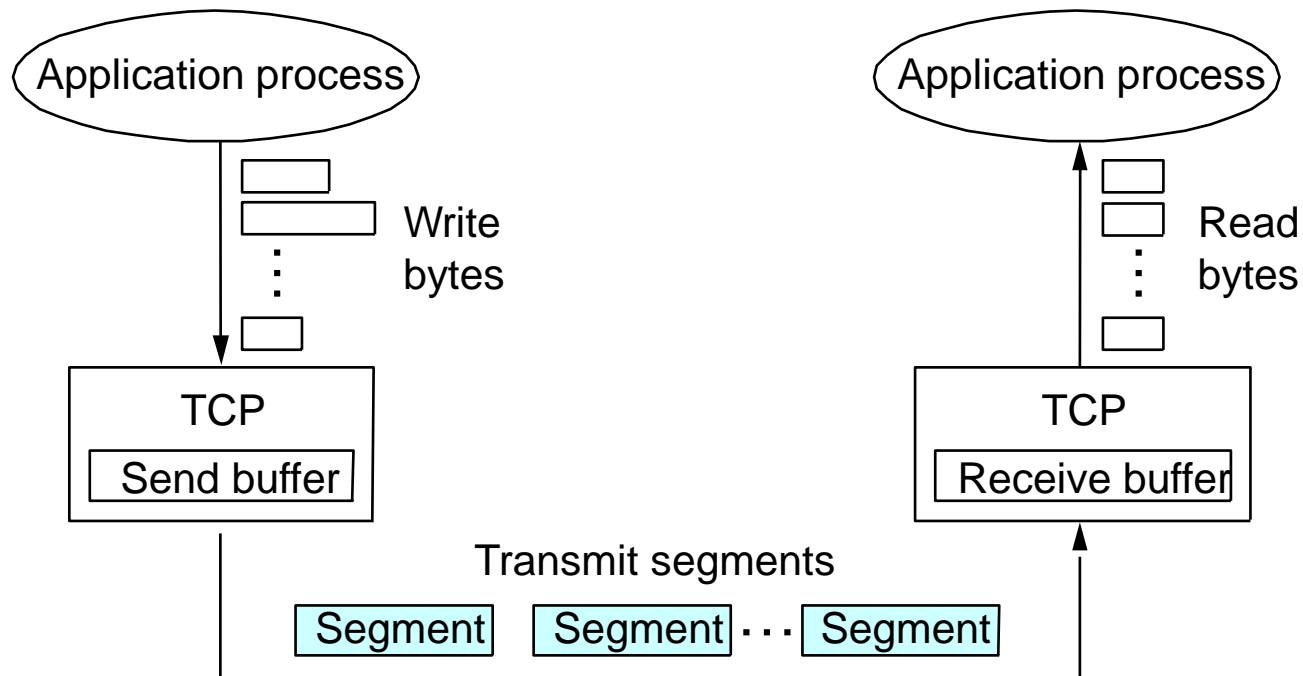
- ▣ Streaming media (e.g., live video)
- ▣ DNS (Domain Name Service)
- ▣ NTP (Network Time Protocol) (synchronizing clocks)
- ▣ FPS multi-player video games (e.g., Call of Duty)
- ▣ Why might UDP be appropriate for these?

Transmission Control Protocol



- Reliable bi-directional **bytestream** between processes
 - ◆ Uses a sliding window protocol for efficient transfer
- Connection-oriented
 - ◆ Conversation between two endpoints with beginning and end
- Flow control (generalization of sliding window)
 - ◆ Prevents sender from over-running receiver buffers
 - ◆ (tell sender how much buffer is left at receiver)
- Congestion control (next lecture)
 - ◆ Prevents sender from over-running network capacity

TCP Delivery





TCP like reliable data link

- Remember we said that when we did reliable data links that TCP would be similar (but end-to-end)
- This is where we “cash in” for all the hard work we did in sliding windows, go back N, Restart Protocols etc
- As a first approximation, TCP takes the bytes the user writes to the queue, packages them in segments, adds a sequence number and does Go Back N
- But there are differences we need to understand

Differences between Data Link reliability and TCP – Part 1



- Network instead of single FIFO link **We'll do this second**
 - packets can be delayed for large amounts of time
 - duplicates can be created by packet looping: delayed duplicates imply need for large sequence numbers.
 - packets can be reordered by route changes.
- Connection management **We'll do this first**
 - Only done for Data Link when a link crashes or comes up
 - Lots of clients dynamically requesting connections
 - HDLC didn't work: here more at stake, have to do it right.

Differences between Data Link reliability and TCP – Part 2



- Data link only needs speed matching between receiver and sender (flow control). Here we also need speed matching between sender and network (congestion control)
- Transport needs to dynamically round-trip delay to set retransmit timers.

We'll do these next lecture




Part 1: Connection Set up


- Both sender and receiver must be ready before we start to transfer the data
 - ◆ Sender and receiver need to agree on a set of parameters
 - ◆ Most important: sequence number space in each direction
 - ◆ Lots of other parameters: e.g., the Maximum Segment Size

- Handshake protocols: setup state between two oblivious endpoints
 - ◆ Need to deal with **delayed** and **reordered** packets
 - ◆ Lets illustrate the problem with the sad tale of John and Martha

Going to the Cafe Protocol

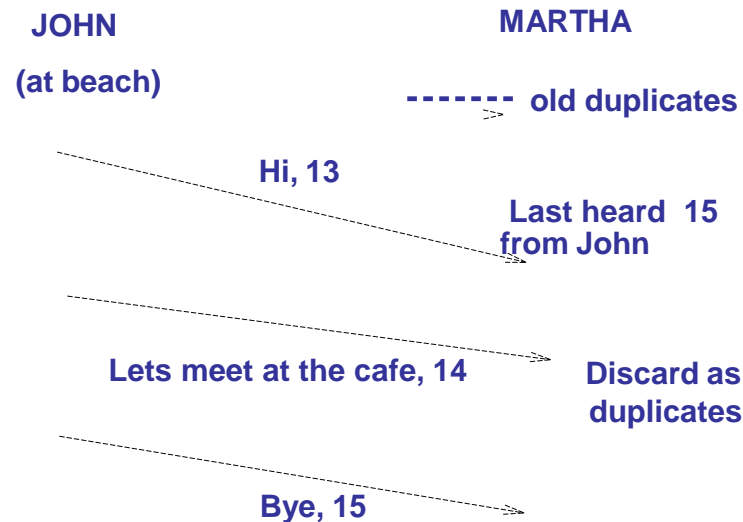


Misery caused by duplicates



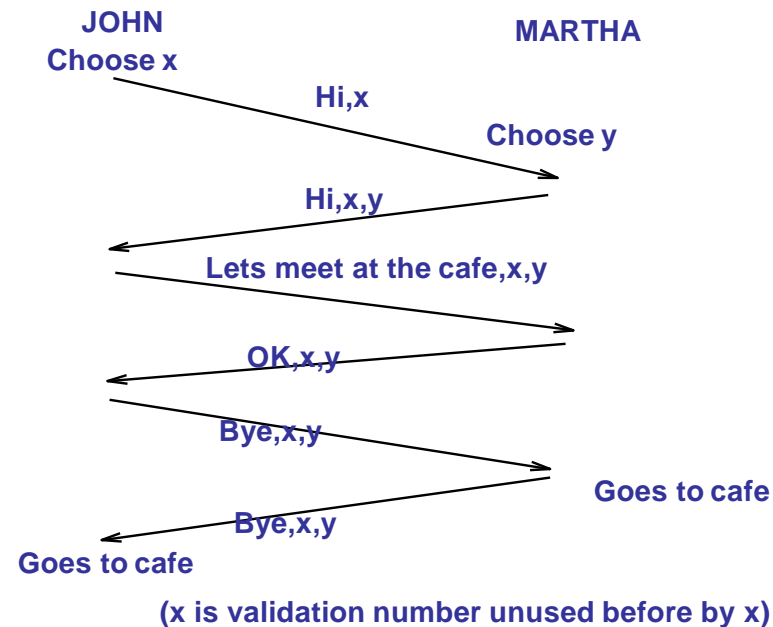


One way out: lots of state




Martha has to remember last number from John for worst-case packet lifetime. Called Timer-based connection management.

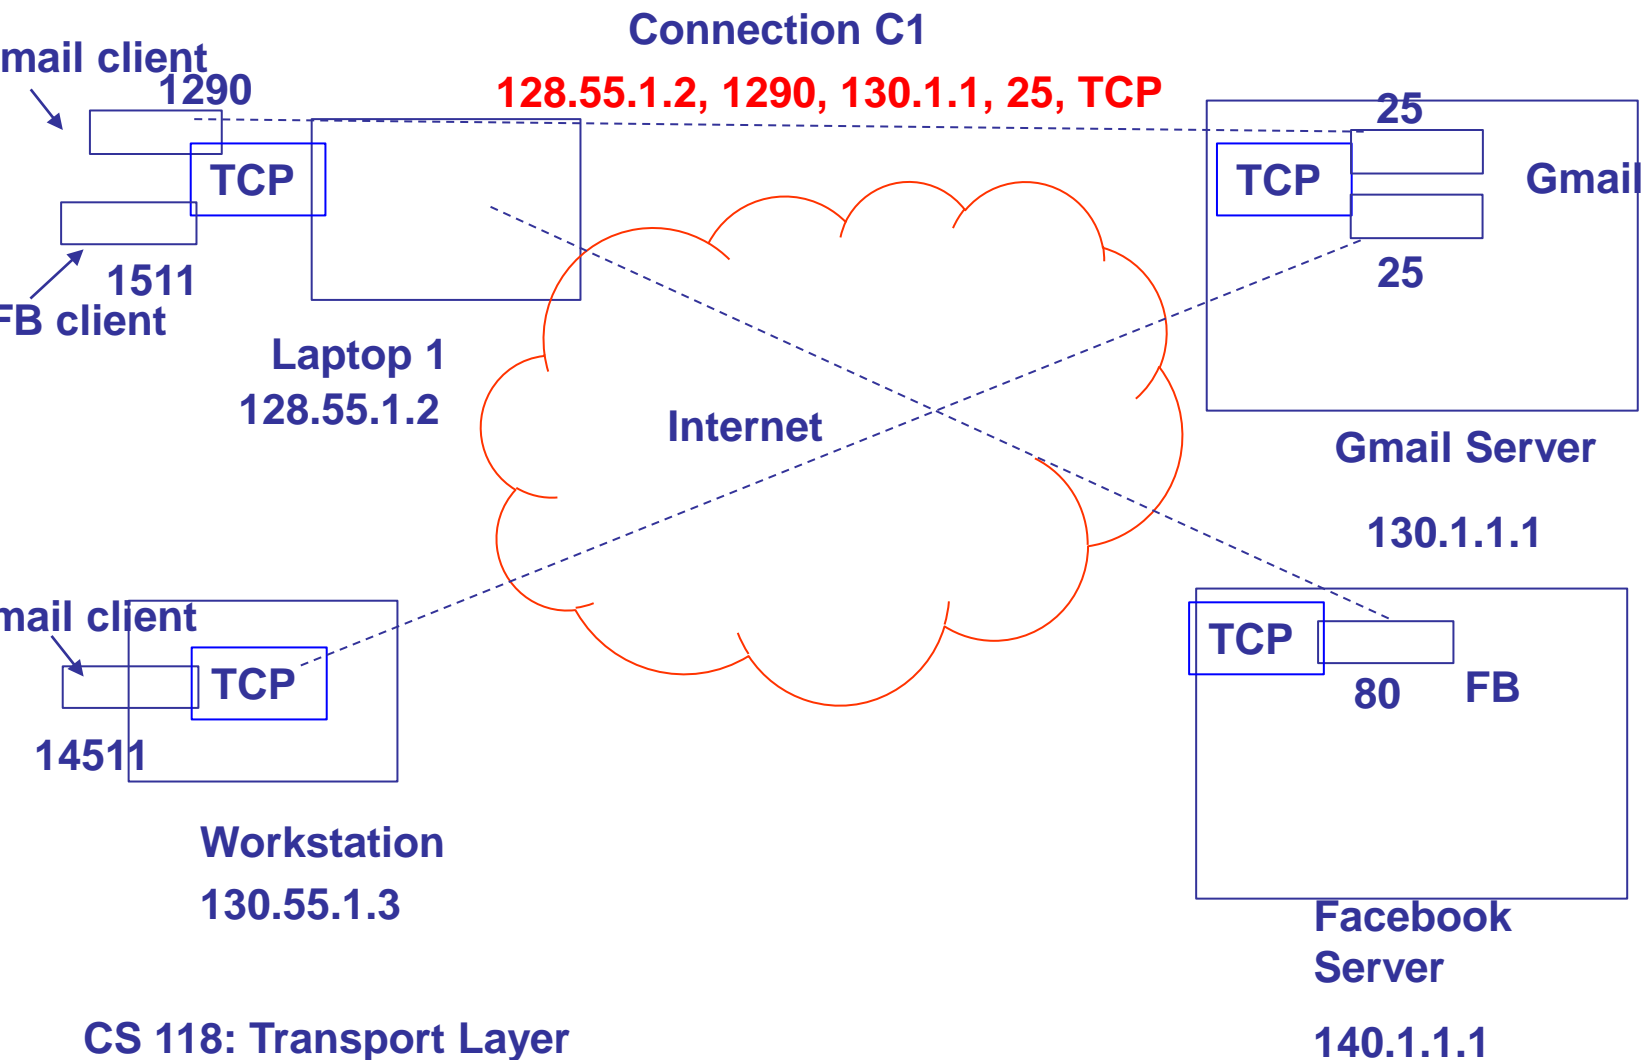
TCP's way: 3-way handshake



Why nonces defend against delayed duplicates



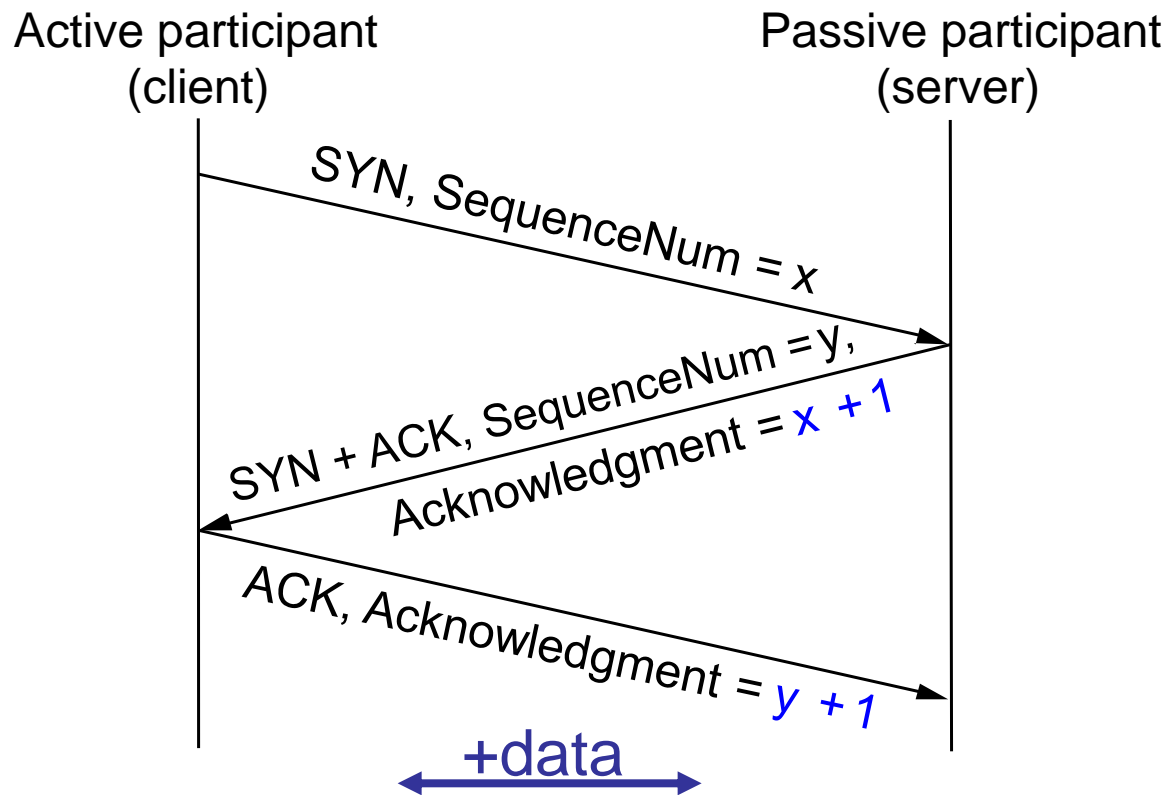
Connection Names: 4-tuples



Three-Way Handshake in TCP



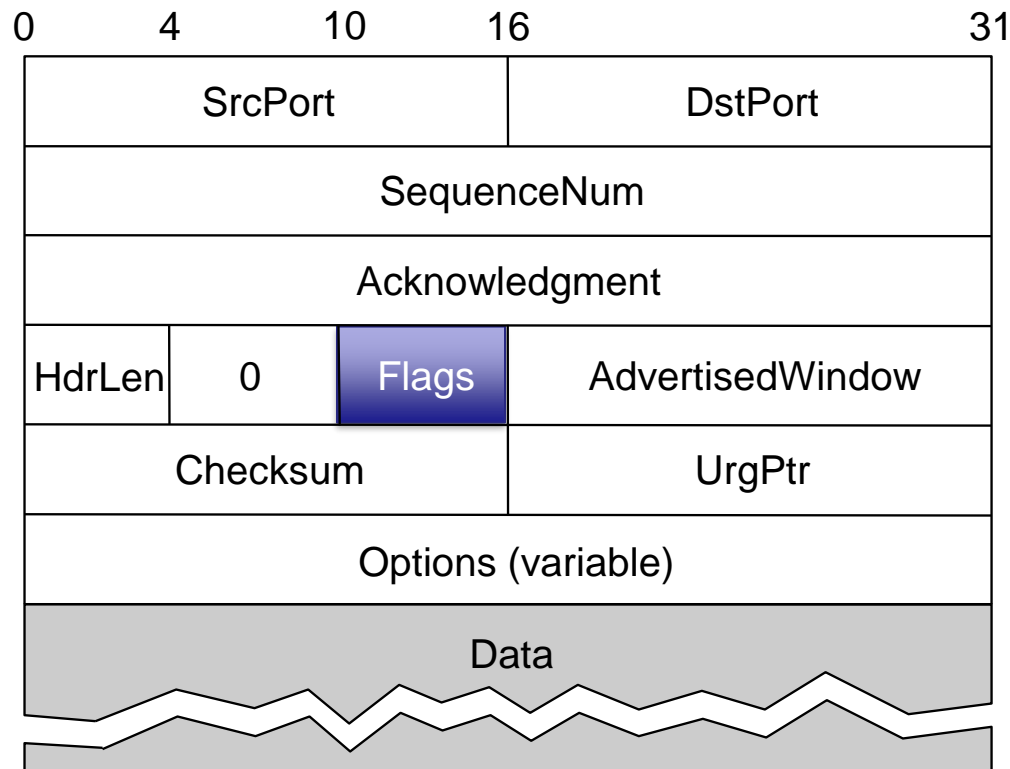
- Opens both directions for transfer





TCP Header Format

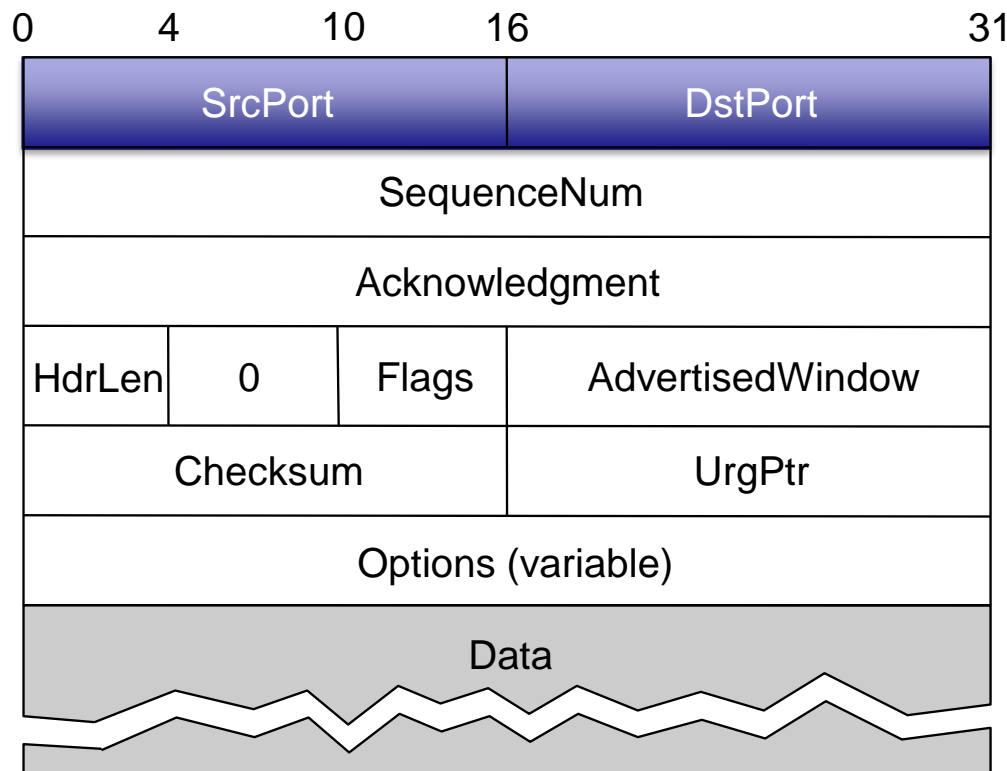
- Flags may be ACK, SYN, FIN, URG, PSH, RST





TCP Header Format

- Ports plus IP addresses identify a connection (4-tuple)





3-way handshake details

- We could abbreviate this setup, but it was chosen to be robust, especially against delayed duplicates
 - ◆ Three-way handshake first described in Tomlinson 1975
- Choice of changing initial sequence numbers (ISNs) minimizes the chance of hosts that crash getting confused by a previous incarnation of a connection
- How to choose ISNs?
 - ◆ Maximize period between reuse
 - ◆ Minimize ability to guess (why?)



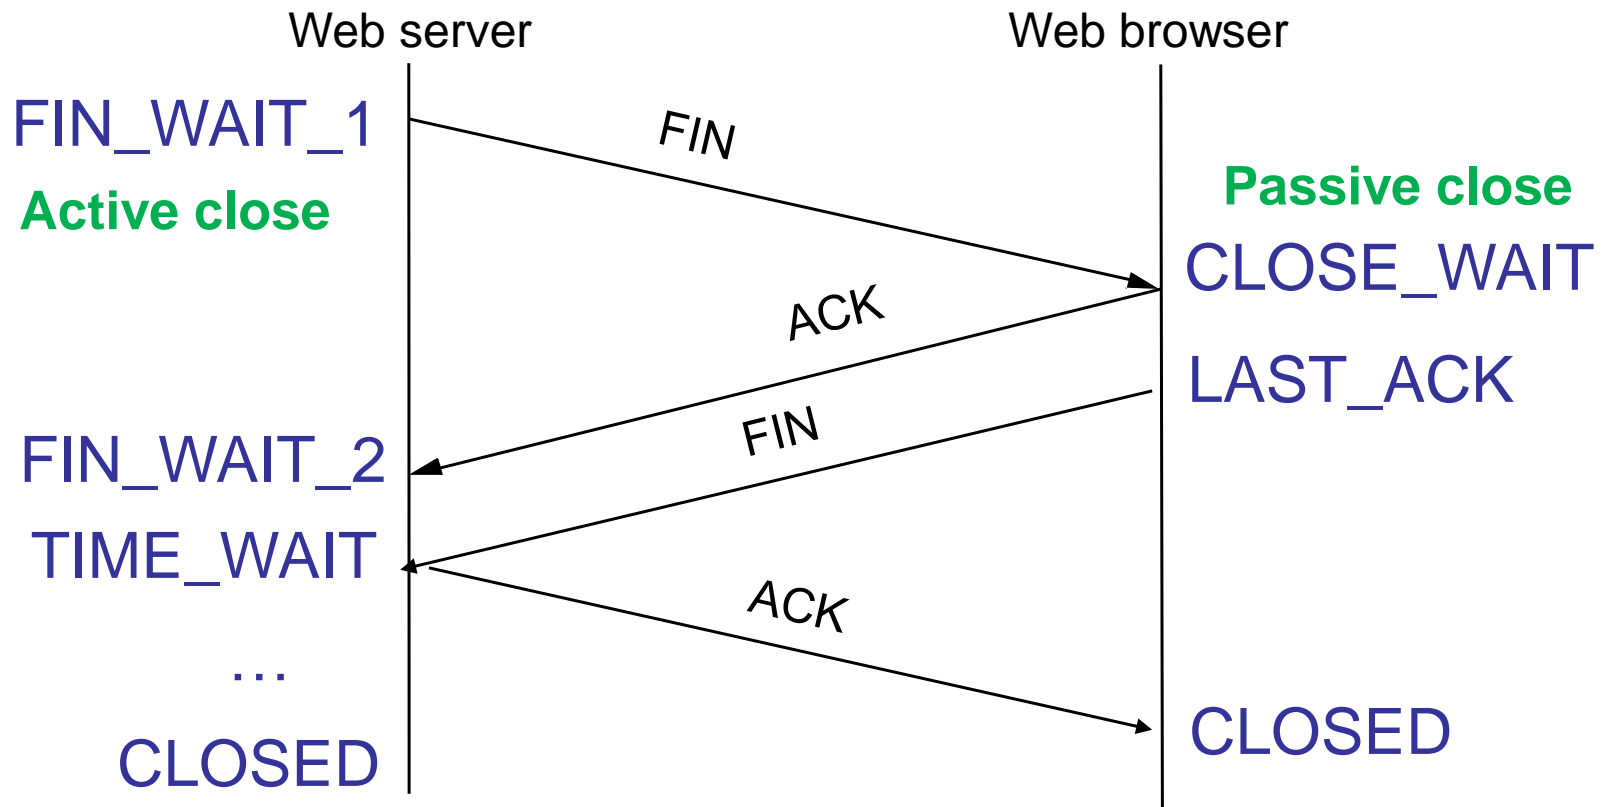
3-way handshake in TCP

- Server: If in LISTEN and SYN arrives, then transition to SYN_RCVD state, replying with ACK+SYN.
- Client: active open, send SYN segment and transition to SYN_SENT.
- Arrival of SYN+ACK causes the client to move to ESTABLISHED and send an ack
- When this ACK arrives the server finally moves to the ESTABLISHED state.

So now how do we disconnect

- 1) Need timers anyway to get rid of connection state to dead nodes.
- 2) However, timer should be large so that "keepalive" hello overhead is low.
- 3) If communication is working, would prefer graceful closing (so receiver process knows quickly) to long timers.
- 4) Hence 3 phase disconnect handshake
After sending disconnect and receiving disconnect ack, both sender and receiver set short timers.

TCP Connection Teardown





The TIME_WAIT State

- We wait $2*MSL$ (maximum segment lifetime of 60 seconds) before completing the close
 - ◆ Why?
- ACK might have been lost and so FIN will be resent
 - ◆ Could interfere with a subsequent connection
- Real life: Abortive close
 - ◆ Don't wait for $2*MSL$, simply send Reset packet (RST)



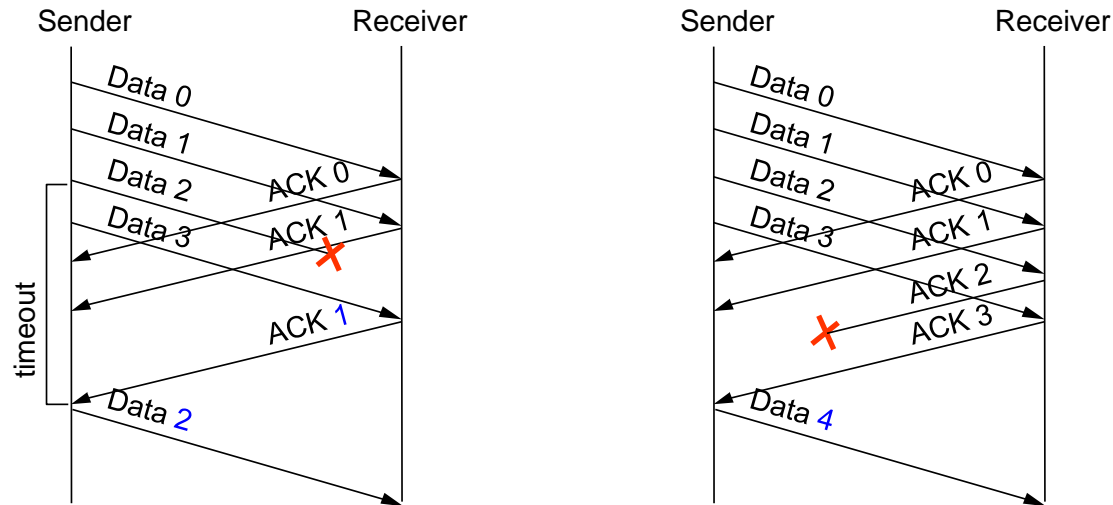
Part 2: Reliable delivery

- Usual sequence numbers except:
 - ◆ Very large to deal with out of order (modulus $> 2W$ etc. only works on FIFO links)
 - ◆ TCP numbers bytes not segments: allows it to change packet size in the middle of a connection
 - ◆ The sequence numbers don't start with 0 but with an ISN.

- Reliable Mechanisms similar except:
 - ◆ TCP has a quicker way to react to lost messages
 - ◆ TCP does a crude form of selective reject not go-back N
 - ◆ TCP does flow control by allowing a dynamic window which receiver can set to reduce traffic rate (next lecture)



Remember Go-Back-N

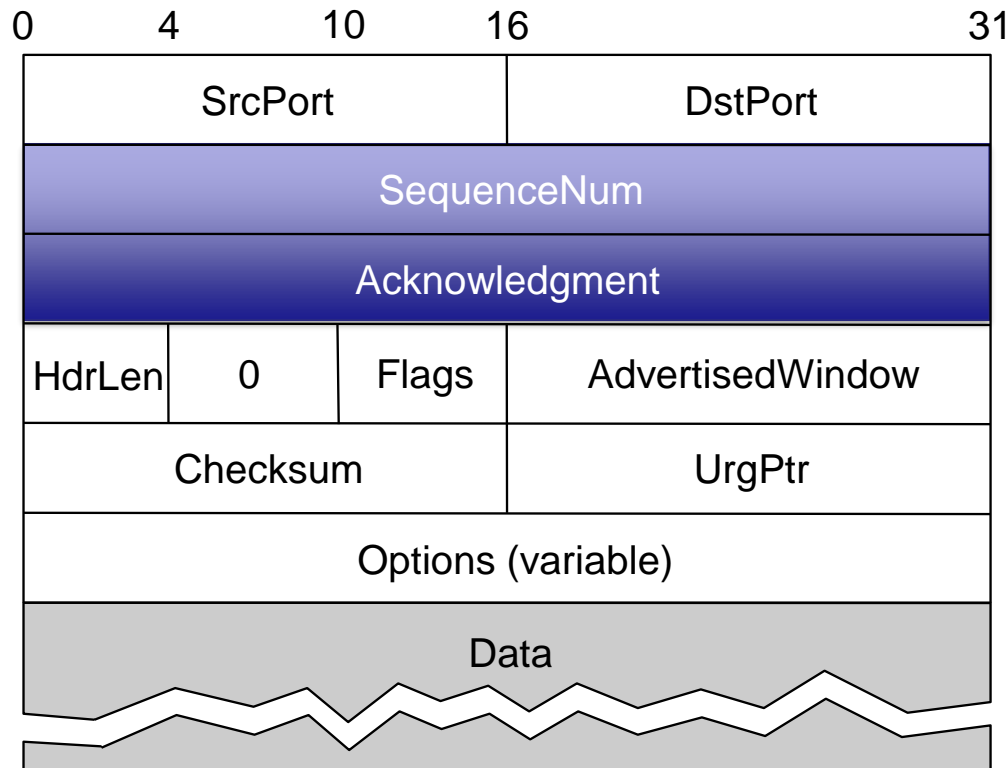


- Retransmit all packets from point of loss
 - ◆ Packets sent after loss event are ignored (i.e., sent again)
- Simple to implement (receiver very simple)
- Sender controls how much data is “in flight”



TCP Header Format

- Sequence, Ack numbers used for the sliding window
 - How big a window? Flow control/congestion control determine



Deciding When to Retransmit

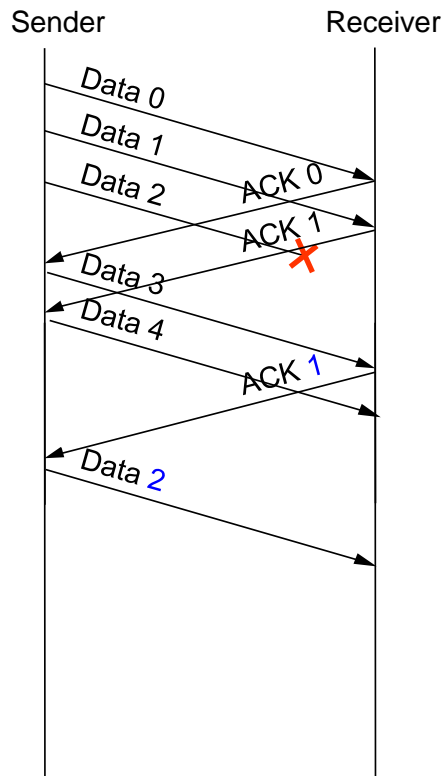


- How do you know when a packet has been lost?
 - ◆ Ultimately sender uses timers to decide when to retransmit

- But how long should the timer be?
 - ◆ Too long: inefficient (large delays, poor use of bandwidth)
 - ◆ Too short: may retransmit unnecessarily (causing extra traffic)


- Right timer is based on the **round-trip time** (RTT)
 - ◆ Which can vary greatly so we need to measure (next lecture)
 - ◆ But OS timer granularity makes it large (msec)
 - ◆ So we need another trick for common case error recovery

TCP Trick: Fast retransmit



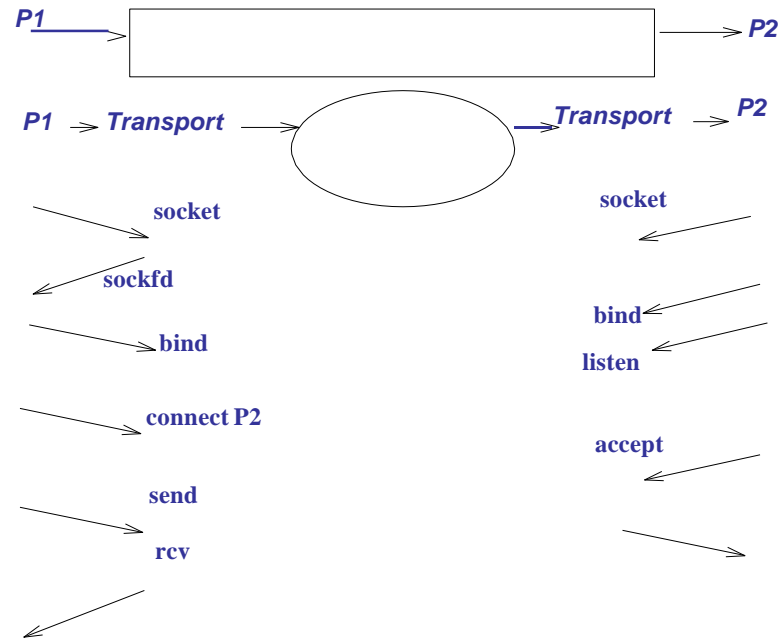
- Don't bother waiting
 - ◆ Receipt of duplicate acknowledgement (**dupACK**) indicates loss
 - ◆ Retransmit immediately
- Used in TCP
 - ◆ Need to be careful if frames can be reordered
 - ◆ Today's TCP identifies a loss if there are **three** duplicate ACKs in a row

Now playing: TCP SACK

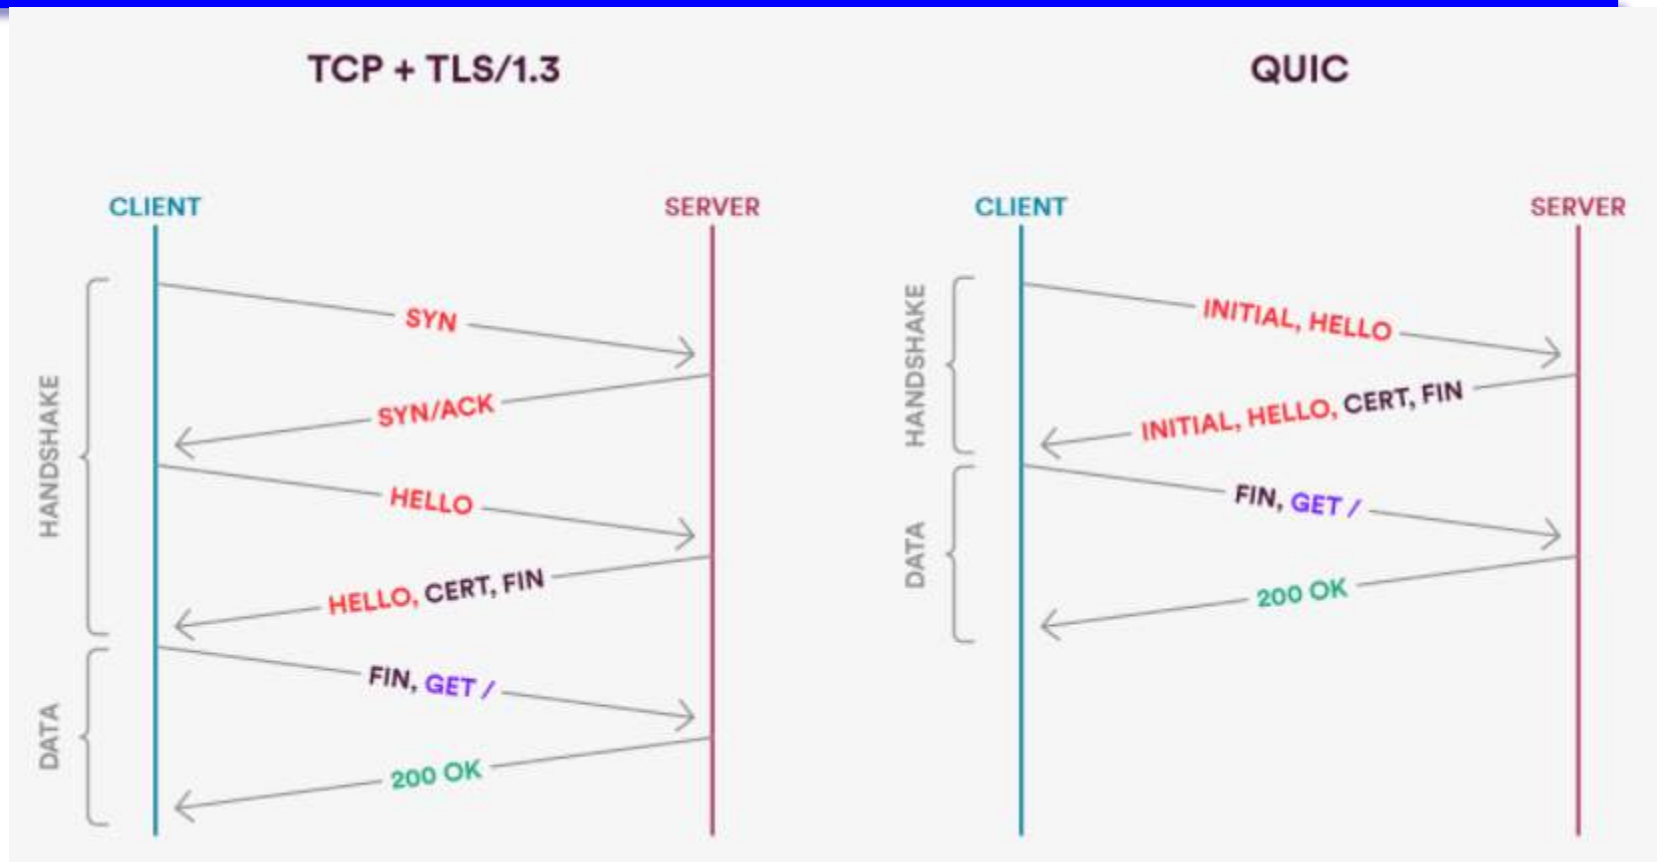


**Has some limitations:
like 3 SACK blocks**

Socket Interface



QUIC first connection to server



QUIC: what's the latency trick



- Idea 1: Combine security handshake and sequence number handshake on first connection to server
- Idea 2: If server remembers information about client, need 0 handshakes on later connections
- Three way handshakes are required because server forgets info on client. Important in old days but no longer as servers have massive memory
- A round trip is a big deal (several hundred msec across US) at today's high speeds



QUIC: what else is new

- Stream multiplexing: multiple streams in a single QUIC connection between client and server for HTTP/2
- No head of line blocking: can do HTTP/2 over a single TCP connection but a single loss stalls all streams. Not so in QUIC
- Shared congestion information: as we will see it takes TCP a long time to ramp up. In QUIC all congestion information is shared.
- Wave of future: 4% of all websites use QUIC (3/2020)