

Layering

- Layers: Application => Transport (TCP) => Routing (IP) => Data Link (Ethernet) => Physical
- Layer numbering starts at 1 for physical
- Protocols are the rules governing horizontal communication
- Interfaces are the rules governing vertical communication
- PDU => horizontal communication
 - TPDU (segment), NPDU (packet), DPDU (frame)
 - N-PDU is an N-SDU (Service Data Unit) together with a layer N header
 - An N-PDU is normally an N-SDU with a layer N header, usually the exact same as the SDU passed to layer N - 1
- Strict layering states that each layer should only look at its header and interface data to do its job

Physical Layer

- Physical Sublayers: transmission, media dependent, coding/decoding

Transmission Sublayer

- The two facets of channel distortion are inertia (lack of bandwidth) and noise
- The Nyquist limit states that we cannot send symbols faster than $2 \cdot \text{bandwidth}$ w/o ISI
- The rate of sending symbols is the signaling/ baud rate
 - The bit rate is the baud rate times # of bits/symbol
- Shannon limit tells us we cannot send faster than $2 \log(S/2N)$ bits/second
 - S = maximum signal strength, N = max noise amp

Clock Recovery/Coding Sublayer

- Compare codes with: guaranteed transitions, transmission efficiency, signal:noise ratio, DC balance, implementation complexity
- Phase locked loops use *many* transitions to generate a clock that is compared to the actual receiver clock => noise will not significantly impact sampling times

Media Sublayer

- Wired media require right of way to install
 - Fiber is light and easy to install and has excellent electrical noise immunity
- Broadcast media (coax/satellite) make possible cheap broadcast protocols, but are insecure
 - Get around need for infrastructure or right of way
- Microwaves are absorbed by rain and infrared is easily absorbed by obstacles
- Low power radio waves and infrared don't require wires or much power

| | Bandwidth | Span | Disadv | Adv |
|--------------|--------------------|-----------|---------------------------|--|
| Twisted Pair | < 10Gbps | 1-2km | low speed | cheap, easy to install |
| Digital Coax | 10-100Mbps | 1-2km | hard to rep, install | broadcast |
| Analog Coax | 100-1000Mbps | 100km | exp. analog amplifiers | cable companies use it now! |
| Satellite | 100-500Mbps | worldwide | prop delay, antennas | no right-of-way issues, cost does not depend on distance |
| Microwave | 10-100Mbps | 100 km | fog outages | no right-of-way |
| Fiber | terabits | 100km | no broadcast, no mobility | security, isolation, bandwidth |
| Infrared RF | < 4 Mbps, 110 Mbps | 3 m, 1 km | obstacles, fog, infrared | wireless |

Figure 12: Pros and Cons of Various Media: a summary

Data Link Layer

- 2 kinds of errors: random bit and burst
- Difference between point-to-point and broadcast links
- PTP: Bits => Framing => Error Detection => Error Recovery (O)
 - Error recovery is historical, when it made sense to transmit every hop
- BC: Bits => Framing => Error Detection => Media Access => Multiplexing (O)
- Accepts packets from the network layer, adds DL header => frame
- E2E argument: only worthwhile guarantees are E2E guarantees
 - DL acks are an optimization, transport acks are the only guarantee
 - Good bit error rate => unreliable DL, bad bit error rate => reliable DL
- Quasi-reliability: undetected error probability and frame loss probability
- E2E argument ignored in practice => slows down performance

Framing

- Needed for timesharing + to create small, manageable units for error recovery/detection
- Flags/Bit Stuffing, Start Flags/Character Count, Start/End Flags from Physical Layer

↳ Stuffing: can't create end flag + can't create flag without stuffing

0111100
011110011110
new flag

1/1 stuffing

01010101
010101010101
flag 5ad

$$\text{Eff.} = \frac{\text{useful bits}}{\text{total bits}}$$
$$\rightarrow \text{Overhead} = \frac{\text{Len(Encoded)} - \text{Len(Orig)}}{\text{Len(Orig)}}$$

Error Detection

- A codeword with Hamming distance d can detect all $d - 1$ bit errors
 - A $2d + 1$ distance code can correct up to d code errors
 - Designer must choose between correction and detection

Error Recovery

- Sent packets must be delivered to the receiver without duplication, loss, or misordering
- Assume: undetected error rate is negligible, FIFO physical layer, frames can be lost, arbitrary delay on physical links
- Stop and Wait:

- Sender has a variable SN (for sender number) initially 0.
- Sender repeats the following loop:

```
1) Accept a new packet from the higher layer if available and store it
   in Buffer B.
2) Transmit a frame (SN, B)
3) If an error-free (ACK, R) frame is received and R is not equal to SN then
   SN = R
   Go to Step 1.
Otherwise if the previous condition does not occur with an arbitrary
timeout period, go to Step 2 after the timeout period.
```

Receiver has a variable RN (for receiver number) initially 0.
Receiver does the following code:

When an error-free data frame (S, D) is received:

```
If S = RN then
    Pass D to higher layer
    RN = RN + 1
```

- Send (Ack, RN)

- There are only two possible consecutive sequences #s in any state

- We can use a single bit \Rightarrow alternating bit

The channel from S to R is either empty or contains a frame M.

The channel from R to S is either empty or contains a frame A.

If a channel contains a frame, then the other channel is empty.

- In doing a proof, you can only use the assumption all invariants hold in the previous state \Rightarrow use invariants to rule out cases

Sliding Window

- Throughput \Rightarrow jobs per second, Latency \Rightarrow time to complete a job
- Transmission rate \Rightarrow rate at which the physical layer sends bits, propagation delay \Rightarrow time it takes a bit to arrive at the receiver
- Stop and wait limits us to 1 frame per round trip delay \Rightarrow problem for all links where transmission rate * propagation delay (bandwidth-delay product/pipe size is large relative to frame size)
- Selective reject is important for large pipe sizes with a large chance of losing frames

Sender code for Go-back N:

Assume all counters are large integers that never wrap.
The sender keeps a lower window L, initially 0.

Send (s,m) (* sender sends or resends s-th data packet *)

The sender can send this frame if and only if:
m corresponds to data packet number s given to sender by client AND
 $L \leq s \leq L + w - 1$ (* only transmit within current window *)

Receive(r, Ack) (* sender absorbs acknowledgement *)

On receipt, sender changes state as follows:
 $L := R$

The receiver keeps an integer R which represents the next sequence number it expects, initially 0.

Receive(s,m) (* receiver gets a data frame *)

On receipt, receiver changes state as follows:
If $s = R$ then (* next frame in sequence *)
 $R := s + 1$
deliver data m to receiver client.

Send(r, Ack) (* we allow receiver to send an ack any time *)
r must be equal to receiver number R at point ack is sent
most implementations send an ack only when a data frame is received

We assume that any unacknowledged frame in current window is periodically resent. In particular, the lowest frame in the current window must be periodically sent to avoid deadlock.

Sender code for selective-reject:

Assume all counters are large integers that never wrap.
The sender keeps a lower window L, initially 0 and a table that indicates which numbers have been acked (this can be optimized to store only a window's worth of such state).

Send (s,m) (* sender sends or resends s-th data packet

The sender can send this frame if and only if:
m corresponds to data packet number s given to sender by client AND
 $L \leq s \leq L + w - 1$ (* only transmit within current window *) AND
s has not been acked.

Receive(R, List, Ack) (* sender absorbs acknowledgement *)

On receipt, sender changes state as follows:

$L = R$

Mark every number in List as being acked in table.

The receiver keeps an integer R which represents the next sequence number it expects, initially 0. The receiver also keeps a table that, for each sequence number, stores a bit indicating whether it has been received and a pointer to the data, if any, being buffered. Once again, this table can be optimized to reduce the amount of storage to be proportional to a window size.

Receive(s,m) (* receiver gets a data frame *)

On receipt, receiver changes state as follows:

If $s > R$ then

Store m in table at position s and set bit in position s
While the bit at position R is not set do
Deliver data at position R
 $R = R + 1$

Send(R, List, Ack) (* we allow receiver to

R must be equal to receiver number R at point ack is sent
List consists of numbers greater than L that have been received.
most implementations send an ack only when a data frame is received

We assume that any unacknowledged frame in current window is periodically resent. In particular, the lowest frame in the current window must be

- Go back N requires a modulus of size $w + 1$, selective reject requires $2w$
- Reliable restarts require either non-volatile memory, probabilistic protocols, or time limits

$E(x)$ is odd-bit error $\Rightarrow E(x)$ has odd # of terms
 $\hookrightarrow E(1) = 1$
 $\hookrightarrow \text{If } E(x)/x + 1 \Rightarrow E(1) = 1 + 1 = 0 \leftarrow$