

CS130: Software Engineering

Lecture 7 Static+Runtime Analysis



<https://forms.gle/GTocHVJ1RuuM4M4TA>

- A word: How's life?
- A tweet: What might you want to know about a program *before* you run it? And why?

Assignment reminders

- Use same note sheet for all assignments
- TL should NOT be submitting any changes
- Project health is important!

Goals of this lecture

- Explain static and runtime analysis
- Discuss the pros and cons of the various approaches
- Show you how to use such analysis in your applications

Key skill: Be able to classify problems as being good candidates for static analysis vs runtime analysis

Static analysis

You've used static analysis before

Compiler warnings (ie, -Wall)

```
$ gcc -Wall foo.cc  
foo.cc:10:2: error: 'my_var' defined but not used [-Werror=unused-variable]
```

Type checking in the compiler

```
$ gcc foo.cc  
foo.cc:5:13: error: cannot convert 'std::string {aka std::basic_string<char>}' to 'int' in initialization
```

Linters

```
$ clang-tidy foo.cc  
foo.cc:12:7: warning: this call will remove at most one item even when multiple items should be removed
```

Compiler warnings

First of all, you should use `-Wall` to enable lots of warnings and use `-Werror` to turn them into errors.

According to `man gcc`, `-Wall` turns on the following:

`-Waddress` `-Warray-bounds` (only with `-O2`) `-Wc++0x-compat` `-Wchar-subscripts` `-Wenum-compare` (in C/ObjC; this is on by default in C++)
`-Wimplicit-int` (C and Objective-C only) `-Wimplicit-function-declaration` (C and Objective-C only) `-Wcomment` `-Wformat` `-Wmain` (only for C/ObjC and unless `-ffreestanding`) `-Wmissing-braces` `-Wnonnull` `-Wparentheses` `-Wpointer-sign` `-Wreorder` `-Wreturn-type`
`-Wsequence-point` `-Wsign-compare` (only in C++) `-Wstrict-aliasing` `-Wstrict-overflow=1` `-Wswitch` `-Wtrigraphs` `-Wuninitialized`
`-Wunknown-pragmas` `-Wunused-function` `-Wunused-label` `-Wunused-value` `-Wunused-variable` `-Wvolatile-register-var`

There are also [tons of other warnings](#) you could enable if you wanted to (eg, `-Wextra`):

`-Wclobbered` `-Wempty-body` `-Wignored-qualifiers` `-Wmissing-field-initializers` `-Wmissing-parameter-type` (C only)
`-Wold-style-declaration` (C only) `-Woverride-init` `-Wsign-compare` `-Wtype-limits` `-Wuninitialized` `-Wunused-parameter`
(only with `-Wunused` or `-Wall`) `-Wunused-but-set-parameter` (only with `-Wunused` or `-Wall`)

Static type checking

Compare:

```
bool Sorted(const std::vector<int>& input) {  
    const int* prev = nullptr;  
    for (const auto& it : input) {  
        if (prev && *prev > it) {  
            return false;  
        }  
        prev = &it;  
    }  
    return true;  
}
```

Immutable

Infer type
when not
important

With:

```
def Sorted(input):  
    prev = None  
    for it in input:  
        if prev and prev > it:  
            return false  
        prev = it  
    return true
```

Return type
unknown

Immutable?

- Static type checking gives you compile time errors about illegal operations
- This is opposed to runtime typed languages that only give you errors at runtime
- You can even do fancy things like type inference (auto in C++x11) at compile time

Linters

```
$ clang-tidy test.cc
```

- Lint was a tool originally developed alongside the C programming language.
- It was originally intended to help catch nonportable constructs.
- Often, you'll find a less pedantic linter built into compiler frontends.

Linters

```
// This is perfectly legal C++. Can you spot the bug?  
{  
    std::lock_guard<std::mutex>(&global_mutex);  
    critical_section();  
}
```

```
// Also legal, but is essentially a 'use-after-free'.  
std::string str = "Hello, world!\n";  
std::vector<std::string> messages;  
messages.emplace_back(std::move(str));  
std::cout << str;
```

- Linters may not seem like static analysis
- But, they correct more than style. Here are some examples:
 - Inaccurate erase/remove
 - Suspicious semicolon
 - Unused RAI
 - Use after move
- Many errors manifest as simple typos that are allowed by the compiler but are likely semantically wrong

What is static analysis?



- A process that inspects the code of your program without executing it directly
- Often will build a control flow graph, though that isn't required
- Looks for patterns in that graph that represent likely problems

What is static analysis?



- As discussed, compiler warnings, type checking and linters are certainly forms of static analysis
- However, most people associate static analysis with a program (other than the compiler) that inspects your code for classes of bugs
- Often this checker is trying to disprove the existence of certain classes bugs in your code

Example: Use after free

```
// Allocate 'a' on the heap
int* a = new int;

// 'a' is still live, so this is OK
*a = 7;

// Free memory associated with 'a'
delete a;

// Best case, this will trigger a SIGSEGV.
// Worst case, this will not trigger a SIGSEGV and silently
// do the wrong thing.
*a = 8;
```

- We can statically determine that 'a' isn't live on the last line
- Value:
 - Helps us avoid SIGSEGVs
 - Helps avoid attacks that may be able to run malicious code by taking advantage of a use-after-free
- Yes, this example is super trivial

Example: Use after free

```
class Foo {  
    public:  
        Foo(int* a) a_(a) { ... }  
        set(int a) { *a_ = a; }  
    private:  
        int* a_;  
};
```

```
Foo build() {  
    int a;  
    return Foo(&a);  
}
```

```
void run() {  
    // Best case, this will trigger a SIGSEGV.  
    // Worst case, this silently do the wrong thing.  
    build().set(7);  
}
```

- ... but this example is more complex
- In this case, `int a` isn't live at the point where it will be used in `run()`
- This is because `Foo` stored a pointer to `int a` as a member
- And then `int a` was freed when it went out of scope at the end of `build()`

Example: Buffer Overrun

```
// Wrong:
int a[10];
memset(a, 0, 100);
// This just stomped on 90*4 bytes past the end of 'a'.
```

```
// Right:
int a[10];
memset(a, 0, sizeof(a));
```

- Reading or writing past the end of a buffer will produce undefined results
- Hopefully you run into a guard page and it causes a SIGSEGV
- Otherwise, it will just silently stomp on memory
- The compiler can often catch this when the size is known at compile time
- Much harder if it is dynamically sized

Example: Buffer Overrun

```
// Wrong:
int a[10];
memset(a, 0, 100);
// This just stomped on 90*4 bytes past the end of 'a'.
```

```
// Right:
int a[10];
memset(a, 0, sizeof(a));
```

```
// Better:
int a[10] = {};
```

- Reading or writing past the end of a buffer will produce undefined results
- Hopefully you run into a guard page and it causes a SIGSEGV
- Otherwise, it will just silently stomp on memory
- The compiler can often catch this when the size is known at compile time
- Much harder if it is dynamically sized

An aside: How are Buffer Overruns Exploited?

```
// Wrong:  
int a[10];  
memset(a, 0, 100);  
// This just stomped on 90*4 bytes past the end of 'a'.
```

```
// Right:  
int a[10];  
memset(a, 0, sizeof(a));
```

```
// More better:  
int a[10] = {};
```

- Can you think of a way to use a buffer overrun to execute malicious code?
- Do you know of anything that can be done at runtime to defend against these sort of attacks?
- If interested in more, [Smashing the Stack for Fun and Profit](#)

Example: Deadlock Detection

```
Mutex mu1, mu2;
```

```
void foo() {  
    mu1.Lock();  
    mu2.Lock();  
    mu2.Unlock();  
    mu1.Unlock();  
}
```

```
void bar() {  
    mu2.Lock();  
    mu1.Lock();  
    mu1.Unlock();  
    mu2.Unlock();  
}
```

```
// It is not safe to run foo() and bar() concurrently
```

- Deadlock is when you have 2+ routines waiting on each other in a cycle.
- Order resource acquisition is one of a few ways to avoid deadlock (see also the famous [dining philosophers problem](#))
- Static analyzers can detect if there is a reachable state where both `foo()` and `bar()` are both concurrently executing and waiting on each other

An aside: Lock Annotation

```
Mutex mu1, mu2;
int a GUARDED_BY(mu1);
int b GUARDED_BY(mu2);

void foo() REQUIRES(mu1, mu2) {
    a = 0;
    b = 0;
}

void test() {
    mu1.Lock();
    foo();           // Warning! Requires mu2.
    mu1.Unlock();
}
```

- You can sometimes annotate your code to help the static analyzer better understand the intended behavior
- In the case of locks, you can explain what mutex guard which vars and then the static analyzer can check that invariant
- More about this in the threading lecture

Example: Uninitialized variable

```
// Wrong:  
int a;  
printf("%d\n", a);  
// This printed random garbage from the stack.
```

```
// Right:  
int a = 0;  
printf("%d\n", a);
```

- Uninitialized vars can result in undefined behavior (the contents of that memory isn't well defined).
- Assign-before-use is often easily identified by compilers or other static checks

Aside: Type checking printf()

```
// Wrong:  
int a;  
printf("%d\n", a);  
// This printed random garbage from the stack.
```

```
// Right:  
int a = 0;  
printf("%d\n", a);
```

- It is kinda crazy that the compiler will type check `printf()` for you.
- For one thing, it is a vararg function.
- It is part of `stdlib`, but checked by the compiler, which is at a different layer of abstraction.
- The compiler implementation needs to mirror `stdlib` implementation, which is annoying to keep in sync.
- That said, it is totally worth it because people always get it wrong!

Example: Dead code

```
const int kConstant = 7;

int foo = 5;
if (foo > kConstant) {
    // All this code is dead...
    [...]
}
```

- It is possible to prove that certain blocks are never reachable
- Here is a contrived example, but you can imagine this being arbitrarily complex
- In general, you can attempt to determine if a guard will always be false

How static analysis works

```
DEFINE DOESITHALT(PROGRAM):  
{  
    RETURN TRUE;  
}
```

THE BIG PICTURE SOLUTION
TO THE HALTING PROBLEM

- In general, static analysis can be reduced to the halting problem, and therefore is undecidable in general.
- That is, the halting is a particular program property that you might want to compute statically, so by simple reduction the halting problem is a static analysis problem. Therefore, you can't compute all properties statically.
- But
 - We can produce approximates.
 - We can require assumptions or annotations to assist

How static analysis works

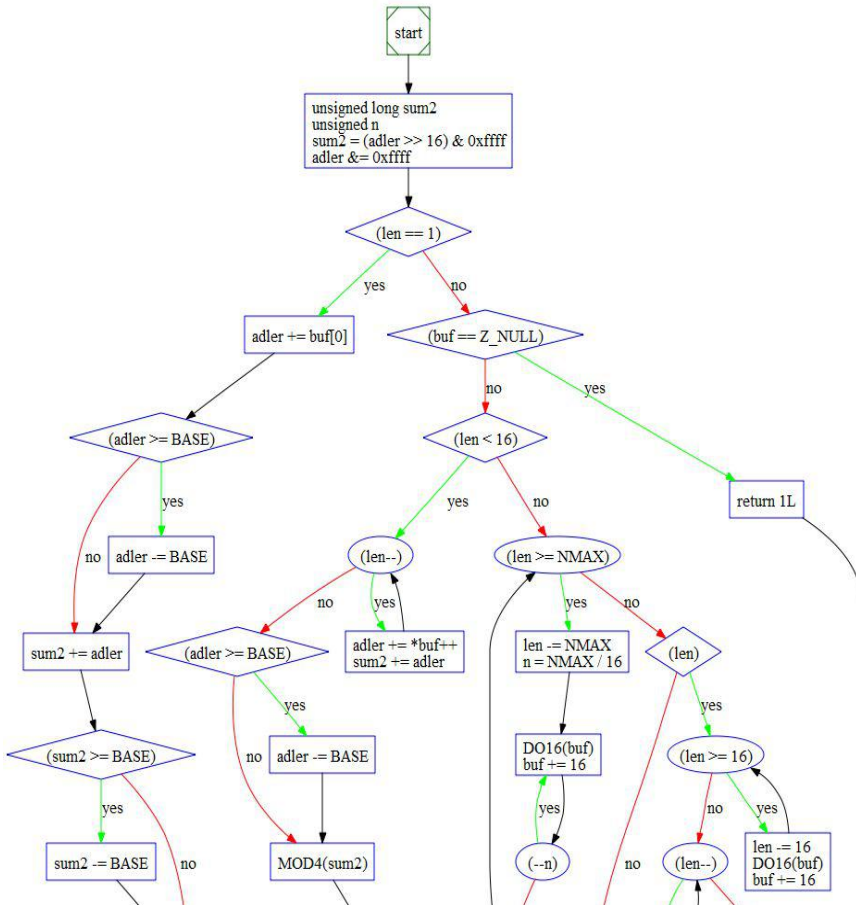
```
DEFINE DOESITHALT(PROGRAM):  
{  
    RETURN TRUE;  
}
```

THE BIG PICTURE SOLUTION
TO THE HALTING PROBLEM

Typical approaches:


- **Abstract interpretation:** model effect of statements on an abstract machine to identify mistakes
- **Data flow analysis:** attempt to determine possible input values based on the control flow graph (similar to general type inference in languages like ocaml)

Abstract Interpretation




- Walk the control flow graph
- Keep track of things that are true when a given unit of code executes
- Determine if invariants are broken
- Example bugs you can catch:
 - Use after free
 - Uninitialized vars
 - Deadlock (if include concurrent execution)

Data Flow Analysis



- Keep track of the set of possible values a variable can take at a given point in the program
- Identify statements that break invariants for a possible value a variable could take on that point
- Example bugs you can catch:
 - buffer overrun
 - dead/unreachable code

Limitations of static analysis



- There are an exponential number of paths through a program
 - Keeping track for each requires an exponential amount of memory
- The range of possible inputs isn't limited much throughout the program
 - Results in false positives

Static Analysis: Overcoming limitations

- Since it's hard to prove/disprove things in static analysis, we often have to pick between low recall (too few bugs get caught) and low precision (correct code gets erroneously flagged)
- One solution: Extend the language

Static Analysis: Overcoming limitations

- Example: Java's `@VisibleForTesting` annotation
- In Java, best practices dictate that class methods should be private unless they need to be used outside of that class.
- Static analysis can flag when methods are non-private but also unused outside of the class context

```
// If no one outside of my
// class is using this method,
// generate a compile-time
// warning

public String getId() {
    ...
}
```

Static Analysis: Overcoming limitations

- What do you do if the only place you need access to the method is for unit testing purposes?
- Want to do 2 things:
 - Make the method public for use by the unit test
 - Maintain Java best practices and act like the method is private by ensuring it is not used anywhere else in the program

```
// Generate compile-time  
// warning if this public  
// method is used outside of  
// either this class or unit  
// tests  
  
@VisibleForTesting  
  
public String getId() {  
  
...  
  
}
```

Static Analysis: Overcoming limitations

- Another example: Typescript
- Javascript does not offer the syntax necessary to do static type checking
- What if you want to build an enterprise application using Javascript, but want the safety that static analysis can provide?

```
// p1, p2, and return types
// should be ints

function myFunction(p1, p2) {
    return p1 * p2;
}
```

Static Analysis: Overcoming limitations

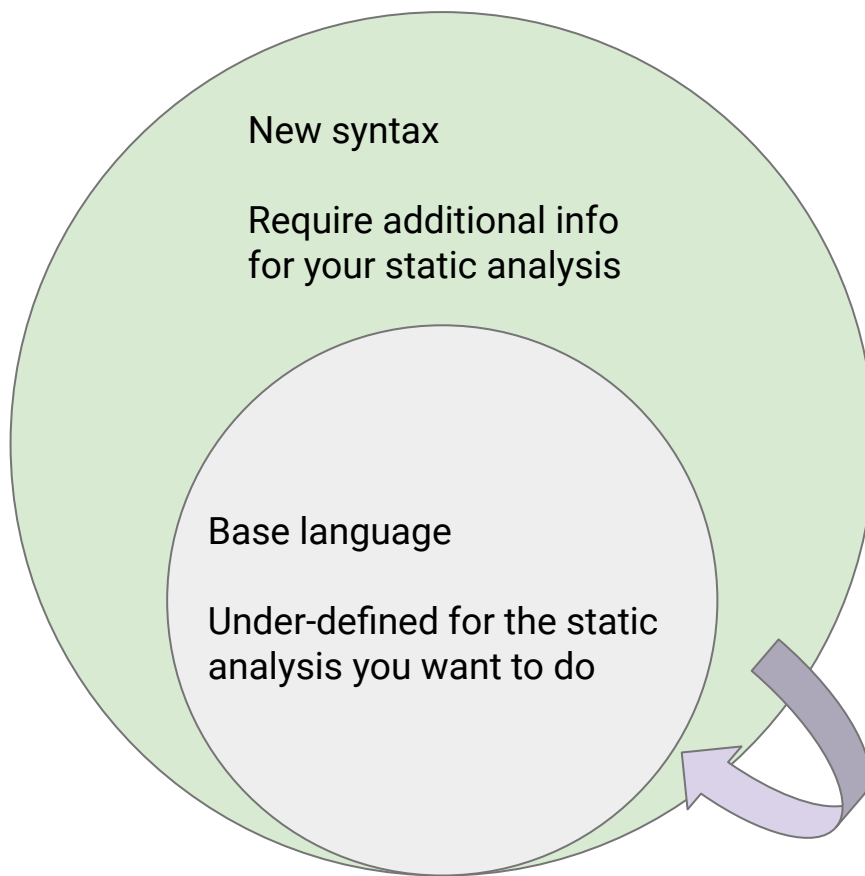
- Solution: Introduce new syntax
- Define “Typescript”, which is a superset of Javascript that can provide detailed typing information
- Then, provide tools:
 - Tool to perform static analysis
 - Tool to automatically generate Javascript from Typescript

In reality, this is the same tool:

The Typescript compiler

```
// p1, p2, and return types
// should be ints

function myFunction(
    p1: int,
    p2: int): int {
    return p1 * p2;
}
```



Tools to convert your new syntax back to base language

Available static analysis tools



- type checking
- -Wall
- lint, clang-tidy
- gcc, clang
- findbugs
- coverity (not free)
- ... and others

clang-format

- Uses LLVM's abstract syntax tree
- Allows automated reformatting of large swaths of code quickly and configurably
- Works with lots of languages
 - Java
 - C++
 - JavaScript
 - Python
 - ObjC

- Used by many large open/closed source projects
 - LLVM
 - Google
 - Chromium
 - Mozilla
 - Apple
 - WebKit
- How do you know this?
 - They implemented their own -styles :)

clang-format

- Doesn't really matter that much when you're the only one reading/writing the code. You might regret messy code later, but will probably still understand it.
- Does matter when **many many** people are all writing code in same project (possibly with differing opinions). Theory: code is read many more times than authored. Also, if your organization has gone through the effort making a style guide, it basically means the formatting matters / you might want to enforce it.
- Ends formatting arguments before they begin. "Just format it."

clang-format

[Live demo](#) ([SkittleParser](#))

FindBugs Demo


FindBugs (1.2.1-dev-20070506) Analysis for jdk1.7.0-b12

Bug Summary	Analysis Information	List bugs by bug category	List bugs by package
-------------	----------------------	---------------------------	----------------------

FindBugs Analysis generated at: Sun, 6 May 2007 03:12:12 -0400

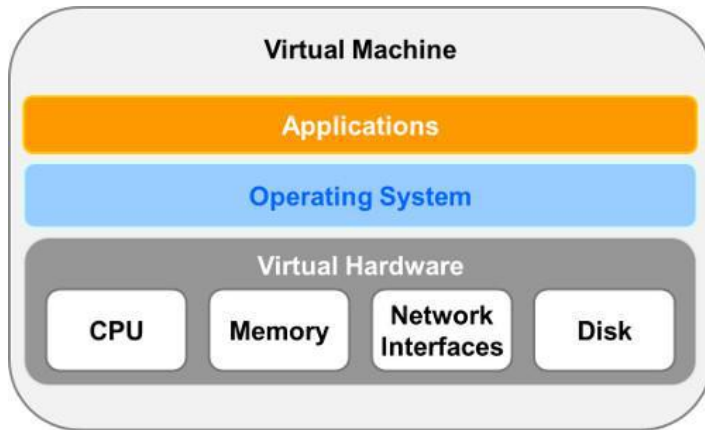
Package	Code Size	Bugs	Bugs p1	Bugs p2	Bugs p3	Bugs Exp.	Ratio
Overall (736 packages), (16445 classes)	963957	3901	259	3642			
com.sun.corba.se.impl.activation	1688	34	5	29			
com.sun.corba.se.impl.copyobject	71	1		1			
com.sun.corba.se.impl.corba	2118	33		33			
com.sun.corba.se.impl.dynamicany	2287	16	3	13			
com.sun.corba.se.impl.encoding	5652	55	1	54			
com.sun.corba.se.impl.interceptors	1979	41		41			
com.sun.corba.se.impl.io	3438	47	2	45			
com.sun.corba.se.impl.ior	1207	14	2	12			
com.sun.corba.se.impl.ior.iiop	457	4		4			
com.sun.corba.se.impl.javax.rmi.CORBA	337	3	1	2			
com.sun.corba.se.impl.logging	9374	8		8			
com.sun.corba.se.impl.naming.cosnaming	799	27	1	26			
com.sun.corba.se.impl.naming.pcossnaming	690	37	4	33			
com.sun.corba.se.impl.oa.poa	2102	31	1	30			
com.sun.corba.se.impl.orb	2324	46	2	44			
com.sun.corba.se.impl.orbutil	3795	25	3	22			
com.sun.corba.se.impl.orbutil.concurrent	320	4		4			
com.sun.corba.se.impl.orbutil.threadpool	357	8		8			
com.sun.corba.se.impl.presentation.rmi	1634	19	2	17			
com.sun.corba.se.impl.protocol	2133	15		15			
com.sun.corba.se.impl.protocol.giopmsgheaders	1861	13	1	12			
com.sun.corba.se.impl.resolver	299	1		1			
com.sun.corba.se.impl.transport	2266	24	1	23			

Runtime analysis



- Alternatively, you could just run the program in an instrumented runtime
- Then, just inspect what happened.
- Think of your brute force debugging sessions where you add `printf()`s until you find the issue.

How runtime analysis works




- Just a virtual machine that is checking for bad states (e.g., SIGSEGV, deadlock, etc.)
- Can keep track of real in-progress state, and pinpoint issues

Limitations of runtime analysis




- Typically, this is slow
- That is, you can't run your code this way in prod (though not all that slow an absolute sense).
- Coverage is limited by test cases

Limitations of runtime analysis



- In some cases, hardware acceleration is available.
- For example, x86 has support for setting a breakpoint when a particular memory address is written to.
- This has virtually no impact on program execution speed.

Available runtime analysis tools



- gcov
- asan
- tsan
- gdb
- valgrind
- memcheck

Demo of gcov + lcov

```
int GreatestOfThree(int a,int b,int c) {
    if ((a > b) && (a > c)){ return a; }
    else if (b > c) { return b; }
    else { return c; }
    return 0;
}

TEST(GreaterTest,AisGreater){
    EXPECT_EQ(3, GreatestOfThree(3,1,2));
};

$ g++ -o main -fprofile-arcs
    -ftest-coverage main.cc
```

```
[=====] Running 3 tests from 1 test case.
[-----] Global test environment set-up.
[-----] 3 tests from GreaterTest
[ RUN      ] GreaterTest.AisGreater
[          OK ] GreaterTest.AisGreater (0 ms)
[ RUN      ] GreaterTest.BisGreater
[          OK ] GreaterTest.BisGreater (0 ms)
[ RUN      ] GreaterTest.CisGreater
[          OK ] GreaterTest.CisGreater (0 ms)
[-----] 3 tests from GreaterTest (0 ms total)
[-----] Global test environment tear-down
[=====] 3 tests from 1 test case ran. (0 ms total)
[ PASSED   ] 3 tests.
```

LCOV - code coverage report

Current view: [top level](#)
Test: [new_coverage.info](#)
Date: 2017-09-26

	Hit	Total	Coverage
Lines:	630	699	90.1 %
Functions:	59	60	98.3 %

Directory	Line Coverage ↕	Functions ↕
	<div></div> 90.1 % 630 / 699	98.3 % 59 / 60

Generated by: [LCOV version 1.10](#)

Demo Video of gdb + Valgrind

```
$ valgrind ls
==211556== Memcheck, a memory error detector
==211556== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==211556== Using Valgrind-3.13.0 and LibVEX; rerun with -h for copyright info
==211556== Command: ls
==211556==
```

<normal output of ls>

```
==211556==
==211556== HEAP SUMMARY:
==211556==      in use at exit: 23,145 bytes in 18 blocks
==211556==    total heap usage: 55 allocs, 37 frees, 62,396 bytes allocated
==211556==
==211556== LEAK SUMMARY:
==211556==    definitely lost: 0 bytes in 0 blocks
==211556==    indirectly lost: 0 bytes in 0 blocks
==211556==    possibly lost: 0 bytes in 0 blocks
==211556==    still reachable: 23,145 bytes in 18 blocks
==211556==          suppressed: 0 bytes in 0 blocks
==211556== Rerun with --leak-check=full to see details of leaked memory
==211556==
==211556== For counts of detected and suppressed errors, rerun with: -v
==211556== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

```
$ valgrind ls
==211556== Memcheck, a memory error detector
==211556== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==211556== Using Valgrind-3.13.0 and LibVEX; rerun with -h for copyright info
==211556== Command: ls
==211556==
<normal output of ls>
==211556==
==211556== HEAP SUMMARY:
==211556==      in use at exit: 23,145 bytes in 18 blocks
==211556==    total heap usage: 55 allocs, 37 frees, 62,396 bytes allocated
==211556==
==211556== LEAK SUMMARY:
==211556==    definitely lost: 0 bytes in 0 blocks
==211556==    indirectly lost: 0 bytes in 0 blocks
==211556==    possibly lost: 0 bytes in 0 blocks
==211556==    still reachable: 23,145 bytes in 18 blocks
==211556==          suppressed: 0 bytes in 0 blocks
==211556== Rerun with --leak-check=full to see details of leaked memory
==211556==
==211556== For counts of detected and suppressed errors, rerun with: -v
==211556== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

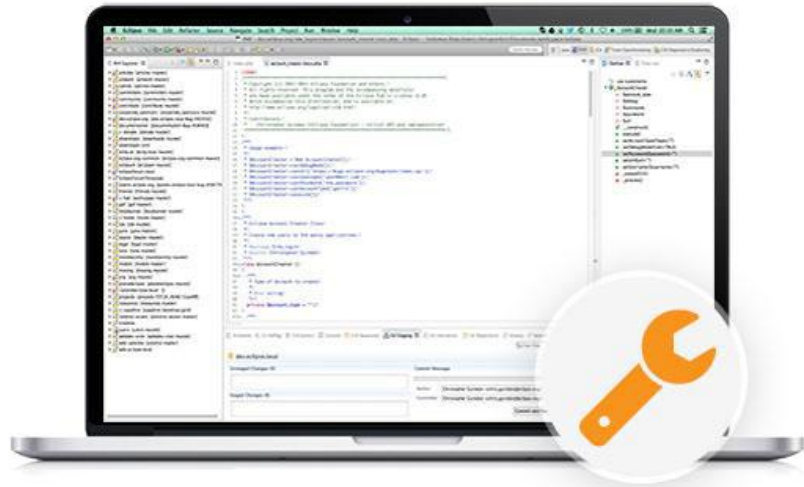

Static Analysis

- **Coverage:** Can cover a somewhat broad class of bugs.
- **Correctness:** Variable false positive rate, depending on bug in question.
- **Issues:** Inherently limited because it is difficult to check for new classes of bugs.

Runtime Analysis

- **Coverage:** Limited by your test coverage.
- **Correctness:** Detects problems that actually happened, not those that might happen.
- **Issues:** Often slow and cannot prove the absence of a bug or class of bugs.

The modern editor



- Modern editors have static analysis built in.
- They essentially have a compiler front end running in the background all the time.
- Can use this analysis to note compile errors (sorta handy).
- More importantly, they can even automate refactors and run static analysis.
- (Note: you can still get this stuff even if you prefer a shell-based editor)

Pitfalls



- Static analysis is no substitute for good coding practice.
- For example:
 - RAI (unique_ptr or scoped_lock) for acquiring resources.
 - Lock acquisition order to avoid deadlock
 - Initialize your variables
 - Code defensively

Let's Check Out:

A word: We can't even solve
The halting problem!
Is static analysis worth it?
A tweet: What tool interested
you the most? Why?

