


CS130: Software Engineering

Lecture 12: Threading and Concurrency

<https://bit.ly/3N1w1SY>

- A word: A random word.
- A tweet: it'll show up after you give the word.



Assignment 7 Notes

Major pieces:

- API request handler for
 - Create
 - Read
 - Update
 - Delete
 - List
- Unit tests
- Integration tests

What you can reuse:

- Common API
- Static Request Handler to read files
- Logic to map from request path to local file path
- Integration test skeleton

Major pieces:

- API request handler for
 - Create
 - Read
 - Update
 - Delete
- Unit tests
- Integration tests

Nice things:

- Create and Update are very similar
 - Create a new file vs write to existing file
- Read is basically static request handler
- Delete is very simple
 - If file exists delete it
- List is simple
 - List the files, return them

Major pieces:

- API request handler for
 - Create
 - Read
 - Update
 - Delete
- Unit tests
- Integration tests

New things:

- Parse data from request for upload
- Abstraction for file operations
- Integration tests with commands other than GET

Possible workflows

1. Parse request into:
 - a. Verb
 - b. Path
 - c. Request data (if present)
2. Request handler methods to:
 - a. Map a request path to a file path
 - b. Read data from a file path
 - c. Write data to a file path
 - d. List file names from file path
 - e. Format file names into JSON text
 - f. Format JSON text into response
3. Create filesystem abstraction for unit tests
 - a. `read_file(path)`
 - b. `write_file(path)`
 - c. `delete_file(path)`
 - d. `list_files(path)`
4. Update integration tests to support POST, PUT, DELETE commands

Concurrency

Background

Concepts you should already know about:


- Threads, processes
- Synchronization primitives: mutexes, semaphores, etc.
- Critical sections, race conditions, deadlocks


Background (refresher)

Concepts you should already know about:

- **Threads, processes**
- Synchronization primitives: mutexes, semaphores, etc.
- Critical sections, race conditions, deadlocks









Background (refresher)

Concepts you should already know about:

- Threads, processes
- **Synchronization primitives: mutexes, semaphores, etc.**
- Critical sections, race conditions, deadlocks





**MAXIMUM
OCCUPANCY
350**




Background (refresher)

Concepts you should already know about:

- Threads, processes
- Synchronization primitives: mutexes, semaphores, etc.
- **Critical sections, race conditions, deadlocks**







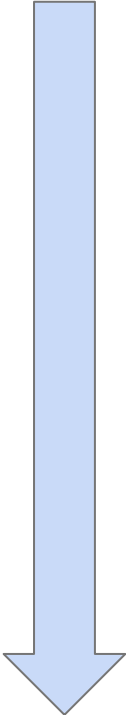


Shared Memory

Threads

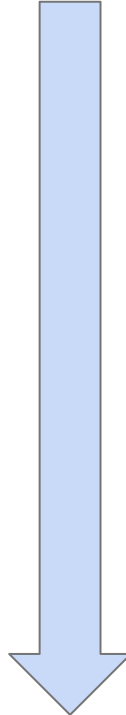


Shared Memory

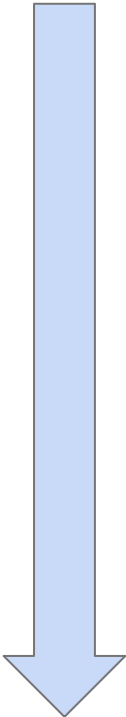


```
1: sharedCounter->nextNumber()  
   // returns 1
```

```
2: sharedCounter->nextNumber()  
   // returns 2?
```

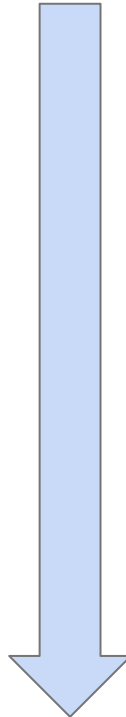


Shared Memory



```
1: sharedCounter->nextNumber()  
  // returns 1
```

```
2: sharedCounter->nextNumber()  
  // returns 2?  
  // oh no returns 3
```



```
1: sharedCounter->nextNumber()  
  // returns 1? Oh no  
  // returns 2 :(
```

```
2: sharedCounter->nextNumber()  
  // returns 2?  
  // not even close, returns 4
```




Why concurrency?

Make things go faster!

- Don't sit around waiting on something. Work on something else
- Modern CPUs have many cores. Keep them busy.
- Make UIs responsive

Concurrency is hard

Concurrency is really hard

Even experienced engineers find it hard to write concurrent code that doesn't race/deadlock/crash.

This lecture is about strategies for coping with concurrency in the real world.

Googlers do it wrong:

Before:

https://chromium.googlesource.com/chromium/src/+/fd5b108c82d56f6022dfbe62a023d1e81ff6f83b/base/files/file_tracing.cc

After:

https://chromium.googlesource.com/chromium/src/+/00dd6010cdda17cecd624c0f274aeb630f31fb83/base/files/file_tracing.cc

Where we messed up:

<https://codereview.chromium.org/2114993003>

Instructors do it wrong:

```
class A-RPCImplementationDetail {  
    String method_; ← local variable  
}
```

```
class B-RPCImpl : A-RPCImpl {  
    void CopyRpc() {  
        rpc_copy.set_string_view(method_);  
    }  
}
```

```
class RpcCopyUser {  
    // somewhere deep in tear-down logic, now using B-RPCImpl.  
    print(rpc.method()); ← accesses const string-view  
}
```

We do it wrong:
www.cs130.org goes down

Where's the problem?

steps:

- name: 'gcr.io/cloud-builders/gcloud-slim'
 entrypoint: 'bash'
 args:
 - '-c'
 - |
 mkdir _site
- name: 'jeekyll/jeekyll:3.8'
 args: ['jeekyll', 'build']
 env:
 - 'JEKYLL_VERSION=3.8'
- name: '18fgsa/html-proofer:latest'
 args: ['_site/', '--disable-external']
- name: 'gcr.io/cloud-builders/gcloud-slim'
 args:
 - 'compute'
 - 'scp'
 - '_site/'
 - 'chronos@\${_GCE_INSTANCE}:\${_DST_PATH}.tmp'
 - '--recurse'
 - '--zone=\${_GCE_ZONE}'
 - '--ssh-key-expire-after=1m'

- name: 'gcr.io/cloud-builders/gcloud-slim'
 args:
 - 'compute'
 - 'ssh'
 - 'chronos@\${_GCE_INSTANCE}'
 - '--zone=\${_GCE_ZONE}'
 - '--ssh-key-expire-after=1m'
 - '--command'
 - |
 if [-d \${_DST_PATH}.last]; then
 rm -rf \${_DST_PATH}.last;
 fi;
 if [-d \${_DST_PATH}]; then
 mv \${_DST_PATH} \${_DST_PATH}.last;
 fi;
 mv \${_DST_PATH}.tmp \${_DST_PATH}

What's the problem?

In one step:

- `scp _site/ X.tmp`

In subsequent step:

- If `X.last` exists, remove it
- If `X` exists, mv it to `X.last`
- Move `X.tmp` to `X`

How can we solve it?

```
- 'compute'
- 'scp'
- '_site/'
- 'chronos@${_GCE_INSTANCE}:${_DST_PATH}.tmp'
```

... then ...

```
if [ -d ${_DST_PATH}.last ]; then
    rm -rf ${_DST_PATH}.last;
fi;
if [ -d ${_DST_PATH} ]; then
    mv ${_DST_PATH} ${_DST_PATH}.last;
fi;
mv ${_DST_PATH}.tmp ${_DST_PATH}
```

The problem:

scp _site/ X.tmp

scp _site/ X.tmp

rm X.last

X -> X.last

X.tmp -> X

rm X.last

X -> X.last

X.tmp -> X (error: X.tmp not found)

The fix:

scp _site/ X.1

scp _site/ X.2

rm X.last

X -> X.last

X.1 -> X

rm X.last

X -> X.last

X.2 -> X

Strategy 1: Be slow

Example: Regression test for ML models

1. Copy models to a staging directory
 - a. Copy baseline models
 - b. Copy test models
2. Evaluate the models on the same inputs. Compare the outputs.

Strategy 1: Be slow

Example: Regression test for ML models

1. Copy models to a staging directory
 - a. Copy baseline models (10 minutes)
 - b. Copy test models (10 minutes)
2. Evaluate the models on the same inputs. Compare the outputs.

Idea: Do 1a and 1b in parallel!

Strategy 1: Be slow

Example: Regression test for ML models

1. Copy models to a staging directory
 - a. Copy baseline models (10 minutes)
 - b. Copy test models (10 minutes)
2. Evaluate the models on the same inputs. Compare the outputs. (2 hours)

~~Idea: Do 1a and 1b in parallel!~~

Lesson: Always profile before you optimize

Strategy 2: Isolate

Example 1: An operating system

- Your first program: `void main() { printf ("Hello, world\n"); }`
- Runs in parallel with other processes, without you having to do anything.

It wasn't always like this.

- In early versions of MacOS, developers had to explicitly yield to the OS.
- Terrible API! Bad for developers (complex), bad for users (crashy).

Strategy 2: Isolate

Example 2: A web indexer

- Input: crawled web pages
- Indexer analyzes each web page, writes what it learned to a database
- Examples of analysis:
 - What other pages does it link to?
 - What entities (e.g. people) are mentioned in the page

```
while (true) {  
    document = GetNextDocument(); // Blocks until next document arrives.  
    ProcessDocument(document);  
}
```

Strategy 2: Isolate

Process each document in a separate thread:

- Concurrency code is isolated to your main loop.
- All the code in **ProcessDocument** is single-threaded.
 - Easier to understand.
 - Other teams may be contributing their own analyses to **ProcessDocument**. Easier for them.
- Natural. When you're processing a document, you shouldn't care about other documents being processed at the same time.
- Throughput more important than latency for a web indexer
 - Care about overall documents processed per second, not time to process each document
 - Just increase the number of threads until your CPU is saturated.

Strategy 2: Isolate

Example 3: A web server

- Run each request in a separate thread

Strategy 3: Use a Library

Example: An RPC server

- We've already isolated each request in its own thread
- But we want to make requests faster

```
void HandleRequest(...) {  
    DoStuff1();  
    DoStuff2();  
}
```

- We want `DoStuff1()` and `DoStuff2()` to happen in parallel.

Threads in C++11

```
#include <thread>
void HandleRequest(...) {
    std::thread t1(DoStuff1);
    std::thread t2(DoStuff2);
    t1.join();
    t2.join();
}
```

Compiler flags: -std=c++11 -pthread

Threads in C++11: Passing Data



Story time

- I wanted to write the simplest possible threading example for this lecture.

Story time

```
#include <iostream>
#include <thread>

void f(int x) {
    std::cout << "f(" << x << ")\n";
}

int main() {
    std::thread t1(f, 1);
    std::thread t2(f, 2);
    t1.join();
    t2.join();
    return 0;
}
```

Story time

```
#include <iostream>
#include <thread>

void f(int x) {
    std::cout << "f(" << x << ")\n";
}

int main() {
    std::thread t1(f, 1);
    std::thread t2(f, 2);
    t1.join();
    t2.join();
    return 0;
}
```

What's the output:

- a) f(1)
f(2)
- b) f(2)
f(1)
- c) Either (a) or (b)

Story time

```
#include <iostream>
#include <thread>

void f(int x) {
    std::cout << "f(" << x << ")\n";
}

int main() {
    std::thread t1(f, 1);
    std::thread t2(f, 2);
    t1.join();
    t2.join();
    return 0;
}
```

What's the output:

- a) f(1)
f(2)
- b) f(2)
f(1)
- c) Either (a) or (b)
- d) f(f(12))
)

Time for a design review

- Find a teammate, describe your problem, ask for advice.
- My experience: 90% of the time, they will suggest a better, simpler, safer way.

An actual design discussion

- Teammate says to me:
"I want to run some code in a thread every 10 seconds."
- Josh already wrote that! (In 2006)
- He also thought about what happens if the code runs for more than 10 seconds.

Threads in C++11: Passing Data

Pretty safe:

- Pass by value. Copy the data.
- `DoStuff()` has its own private copy that no other thread can touch.

```
void DoStuff(int x) {  
    std::cout << x << std::endl;  // 2, for sure.  
}
```

```
int i = 2;  
std::thread t(DoStuff, i);
```

Threads in C++11: Passing Data

Not so safe:

- Pass by const reference or ptr
- `DoStuff()` can't change the data ... but another thread can!

```
void DoStuff(const int& x) {  
    std::cout << x << std::endl;  
}
```

```
int i = 2;  
std::thread t(DoStuff, std::ref(i));
```

```
void DoStuff(const int* x) {  
    std::cout << *x << std::endl;  
}
```

```
int i = 2;  
std::thread t(DoStuff, &i);
```

Threads in C++11: Passing Data

Not so safe:

- Pass by const reference or ptr
- `DoStuff()` can't change the data ... but another thread can!

```
void DoStuff(const int& x) {  
    std::cout << x << std::endl;  
}
```

```
int i = 2;  
std::thread t(DoStuff, std::ref(i));  
i = 3;
```

```
void DoStuff(const int* x) {  
    std::cout << *x << std::endl;  
}
```

```
int i = 2;  
std::thread t(DoStuff, &i);  
i = 3;
```

Threads in C++11: Passing Data

Asking for trouble:

- Pass a (non-const) pointer.
- `DoStuff()` is practically promising to change its value.

```
void DoStuff(int* x) {  
    *x = 4;  
}
```

```
int i = 2;  
std::thread t(DoStuff, &i);  
// What is i now?  
i = 3; // Surely i == 3 now!
```

Threads in C++11: Returning values

```
int DoStuff() { ... }  
std::thread t(DoStuff);  
t.join(); // return value of DoStuff is lost.
```


Threads in C++11: Returning values

```
#include <future>
```

```
int Square(int x) {  
    return x*x;  
}
```

```
auto a = std::async(Square, 2);  
int v = a.get(); // v == 4
```

Threads in C++11: Returning values

```
#include <future>
```

```
int Square(int x) {  
    return x*x;  
}
```

```
std::future<int> a = std::async(Square, 2);  
int v = a.get(); // v == 4
```

Strategy 4: Use a mutex

Example: A thread-safe counter

```
Counter counter;  
counter.Increment("requests", 1);  
counter.Increment("errors", 1);
```

Strategy 4: Use a mutex

Example: A thread-safe counter

```
class Counter {  
    public:  
        void Increment(const string name, int by);  
        void Get(const string name);  
    private:  
        map<string, int> counters_; // STL containers not threadsafe.  
};
```

Strategy 4: Use a mutex

```
void Counter::Increment(const string name, int by) {  
    counter_mutex_.lock();  
    counters_[name] += by;  
    counter_mutex_.unlock();  
}
```

Strategy 4: Use a mutex

Example: A thread-safe counter

```
#include <mutex>

class Counter {
public:
    void Increment(const string name, int by);
    void Get(const string name);
private:
    map<string, int> counters_; // STL containers not threadsafe.
    std::mutex counter_mutex_;
};
```

Strategy 4: Use a mutex

At this point, you may get nervous, and start wondering:

- Is the mutex going to be too slow?
- Do I need to worry about lock contention?
- Should I use a spinlock?
- Should I use a lock-free algorithm?
- Should I use an atomic increment hardware operation?
- Do I need richer semantics? Should I use a semaphore?

Strategy 4: Use a mutex

At this point, you may get nervous, and start wondering:

- Is the mutex going to be too slow? **No!**
- Do I need to worry about lock contention? **No!**
- Should I use a spinlock? **No!**
- Should I use a lock-free algorithm? **No!**
- Should I use an atomic increment hardware operation? **No!**
- Do I need richer semantics? Should I use a semaphore? **No!**

The answer to all these questions is (almost always) "no".

Just use a mutex.

A cautionary example

```
class SomeClass {  
    private:  
        // Helper is expensive to construct, and may not be needed.  
        Helper* helper_ = nullptr;  
        Helper* GetHelper() {  
            if (helper_ == nullptr) helper_ = new Helper;  
            return helper_;  
        }  
};
```

A cautionary example

How can we make `GetHelper()` thread-safe?

- We want to construct `helper_` exactly once, and always return the same value

But then you might think...

How can we make it fast?

- After we've constructed `helper_`, we don't need to lock to return it (faster).
- We only need to lock if `helper_` is null, to avoid constructing it twice.

A cautionary example

Clever ... and wrong:

```
Helper* GetHelper() {  
    if (helper_ == nullptr) { // Fast!  
        helper_mutex_.lock(); // Slow!  
        // GetHelper could have been called while waiting to lock.  
        if (helper_ == nullptr) {  
            helper_ = new Helper;  
        }  
        helper_mutex_.unlock();  
    }  
    return helper_;  
}
```

Correct, simpler, and fast enough

```
Helper* GetHelper() {  
    helper_mutex_.lock();  
    if (helper_ == nullptr) {  
        helper_ = new Helper;  
    }  
    helper_mutex_.unlock();  
    return helper_;  
}
```

Strategy 5: Never forget to unlock

Recall:

```
class Counter {  
    public:  
        void Increment(const string name, int by);  
        void Get(const string name);  
    private:  
        map<string, int> counters_;  
        std::mutex counter_mutex_;  
};
```

Strategy 5: Never forget to unlock

Let's implement `Get()`:

```
int Counter::Get(const string name) {  
    counter_mutex_.lock();  
    return counters_[name];  
    counter_mutex_.unlock(); // Uh-oh  
}
```

Strategy 5: Never forget to unlock

Let's implement `Get()`:

```
int Counter::Get(const string name) {  
    std::lock_guard<std::mutex> lock(counter_mutex_);  
    return counters_[name];  
    // Mutex is unlocked when lock goes out of scope.  
}
```

lock_guard and unique_ptr

```
template<class T>
class lock_guard {
public:
    lock_guard(T& m) : mutex_(m) {
        mutex_.lock();
    }
    ~lock_guard() {
        mutex_.unlock();
    }
private:
    T& mutex_;
};
```

```
template<class T>
class unique_ptr {
public:
    unique_ptr(T* p) : ptr_(p) {}
    ~unique_ptr() {
        delete ptr_;
    }
private:
    T* ptr_;
};
```


Strategy 6: Use tools

Clang thread-safety analysis

- Add annotations which are checked at compile-time
- Developed and used extensively at Google
- Available as of Clang 3.5
- A little tricky to get working. Requires a custom header file.

Strategy 6: Use tools

```
#include "mutex.h" // Custom header. Wraps std::mutex.
class Counter {
public:
    void Increment(const string name, int by);
    void Get(const string name);
private:
    map<string, int> counters_ GUARDED_BY(counter_mutex_);
    Mutex counter_mutex_; // Wrapped type
};
```

Forgetting to lock

```
void Counter::Increment(const string name, int by) {  
    counters_[name] += by;  
}
```

```
$ clang-3.5 -std=c++11 -c -Wthread-safety thread.cc
```

```
thread.cc:8:5: warning: reading variable 'counters_' requires  
holding mutex 'counter_mutex_' [-Wthread-safety-analysis]
```

Forgetting to unlock

```
void Counter::Increment(const string name, int by) {  
    counter_mutex_.lock();  
    counters_[name] += by;  
}
```

```
$ clang-3.5 -std=c++11 -c -Wthread-safety thread.cc
```

```
thread.cc:10:3: warning: mutex 'counter_mutex_' is still held at the  
end of function [-Wthread-safety-analysis]
```

Strategy 7: Comment your code

When writing a new class, add a comment about its thread-safety

`// Thread-safe` → Can be called safely from multiple threads.

`// Thread-compatible` → Caller needs to do their own locking.

`// Thread-hostile` → Unsafe even if you lock it.

Refactoring

What is this “refactoring” you speak of?



- Simply, it is rewriting (editing in the traditional sense) code to improve some property.
- In this case, we are restructuring the code to be more testable.
- Could also refactor to make it more maintainable:
 - Divide up long functions (extract method)
 - Make a class do fewer things (extract class)

Refactoring Examples

Before:

```
void HandleRequest(  
    string url,  
    String body,  
    String request_type,  
    map<string, string> headers,  
    MimeType* mime_type,  
    map<string, string>* headers,  
    string* response_body);
```

After:

```
Reply HandleRequest(Request req);
```

- [Introduce param object](#) when a method's param list gets too long
- Certain functions have many params
 - Often triggered by dependency injection or tunability
- Can create a param object (also known as an options struct) to encapsulate these params
- Nice side effect: you can define defaults and have a bit more control over values before the function executes

Refactoring Examples

Before:

```
...
if url.find(piece) == config_chunk {
    Return my_echo_handler.Handle(request);
}
...
```

After:

```
RequestHandler* h =
CreateHandlerForConfig(config_chunk.name,
config_chunk);

CreateHandlerForConfig(
    const string& name,
    const NginxConfig& config) {
    if (name == "static") {
        return new StaticHandler::Init(config); }
    ...
}
```

- [Extract Function](#) when a method is too long
- Create a helper to encapsulate a bit of logic

Refactoring Examples In Use

```
void ProcessRequest(tcp::socket* sock) {  
    while (true) {  
        const int kBufferLength = 1024;  
        char data[kBufferLength];  
  
        boost::system::error_code error;  
        const size_t length =  
            sock->read_some(boost::asio::buffer(data), error);  
  
        HandleRequest(data, length);  
    }  
}  
  
void HandleRequest(const char* buf, const size_t length) {  
    [...]  
}
```

Replace magic number with symbolic constant

Extract Method

Introduce param object (possibly)

Refactoring Approaches

- Write tests first!
 - Separate no-functionality-change architecture refactoring changes from testing changes.
 - Ensure your tests are functionally-driven (state, not interaction)
 - Introducing new boundaries means new tests -- old tests mean to remain intact
- Updating methods on existing classes means migrating the tests too
 - Frequently easier to do this BEFORE making architectural changes.
 - For small chunks of code, can do this simultaneously.
 - Usually easiest to proceed “outside-in”
 - Make new methods and keep the old, one is gold, the other horrible dead-weight you want to delete ASAP.
 - Implement new methods in terms of old ones and update tests to new API.
 - Migrate implementation separately (hopefully without needing to touch tests).
- Sometimes you have to re-implement
 - If your present code is too distant from the goal, you may need to do a large-scale replacement.

Refactoring Risks



// TODO: finish migrating to new “Fish” interface by the end of 2014

Refactoring Tradeoffs

- You are trying to create value in a refactor.
 - More flexible interfaces
 - Better feature velocity
 - More compatibility
 - Support larger team
 - Better performance
 - ...
- Dealing with refactoring risks
 - Have a well-motivated plan!
 - Thorough milestones
 - Acceptable mid-point states (lots of places to stop and/or pause and have things better than before)
 - Getting to 100% can take a lot of effort.

<https://bit.ly/3vZ5SzU>

Two tweets:

- What's the best part of the course so far?
- What more would you like to see?

