

Performance

CS 130, Lecture 14

Let's check in

<https://forms.gle/vAtgTkuY3LMVS9os7>

A word - What's it like working with other people's code?

A word - What's the program with the worst performance you have ever used?

A tweet - How do you plan out your own work capacity?



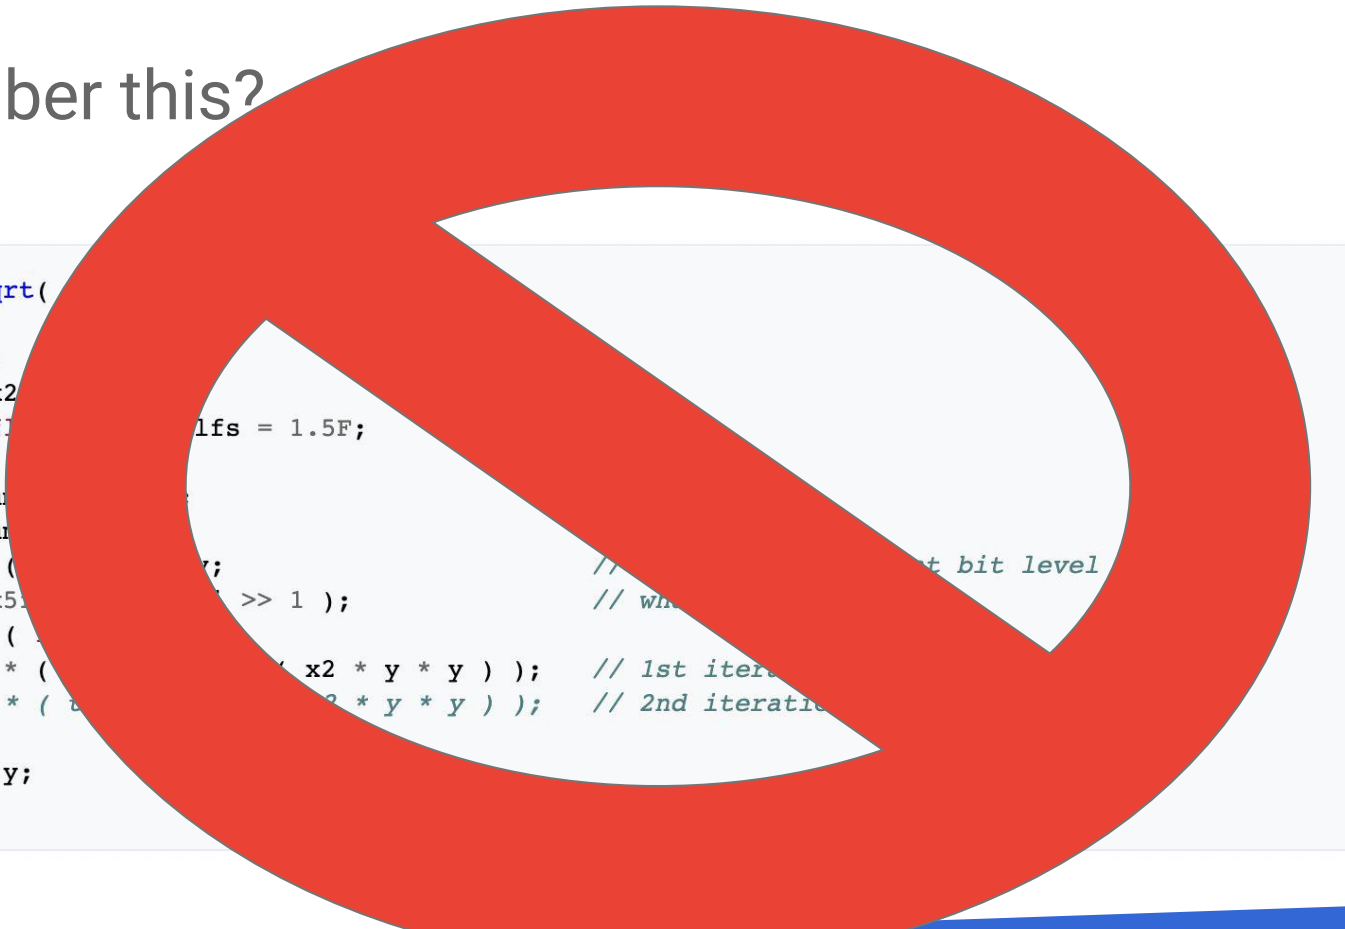
Remember this?

```
float Q_rsqrt( float number )
{
    long i;
    float x2, y;
    const float threehalfs = 1.5F;

    x2 = number * 0.5F;
    y = number;
    i = * ( long * ) &y;           // evil floating point bit level hacking
    i = 0x5f3759df - ( i >> 1 );   // what the fuck?
    y = * ( float * ) &i;
    y = y * ( threehalfs - ( x2 * y * y ) ); // 1st iteration
    // y = y * ( threehalfs - ( x2 * y * y ) ); // 2nd iteration, this can be removed

    return y;
}
```

Remember this?



```
float Q_rsqrt(  
{  
    long i;  
    float x2;  
    const float flfs = 1.5F;  
  
    x2 = num;  
    y = num;  
    i = * ( &y ); // get bit level  
    i = 0x5f3759df >> 1 ); // wh  
    y = * ( &y );  
    y = y * ( 1.5f - 0.5f * x2 * y * y ); // 1st iteration  
    // y = y * ( 1.5f - 0.5f * x2 * y * y ); // 2nd iteration  
  
    return y;  
}
```

Webserver Performance


What are the right performance metrics for a webserver?

Webserver Performance

1. QPS
2. Latency
3. Memory footprint
4. Concurrent connections
5. Bandwidth

~~Webserver Performance~~

- ~~1. QPS~~
- ~~2. Latency~~
- ~~3. Memory footprint~~
- ~~4. Concurrent connections~~
- ~~5. Bandwidth~~



We're tired of talking about web servers

And there's not enough time to optimize
your web servers anyways, sooo...

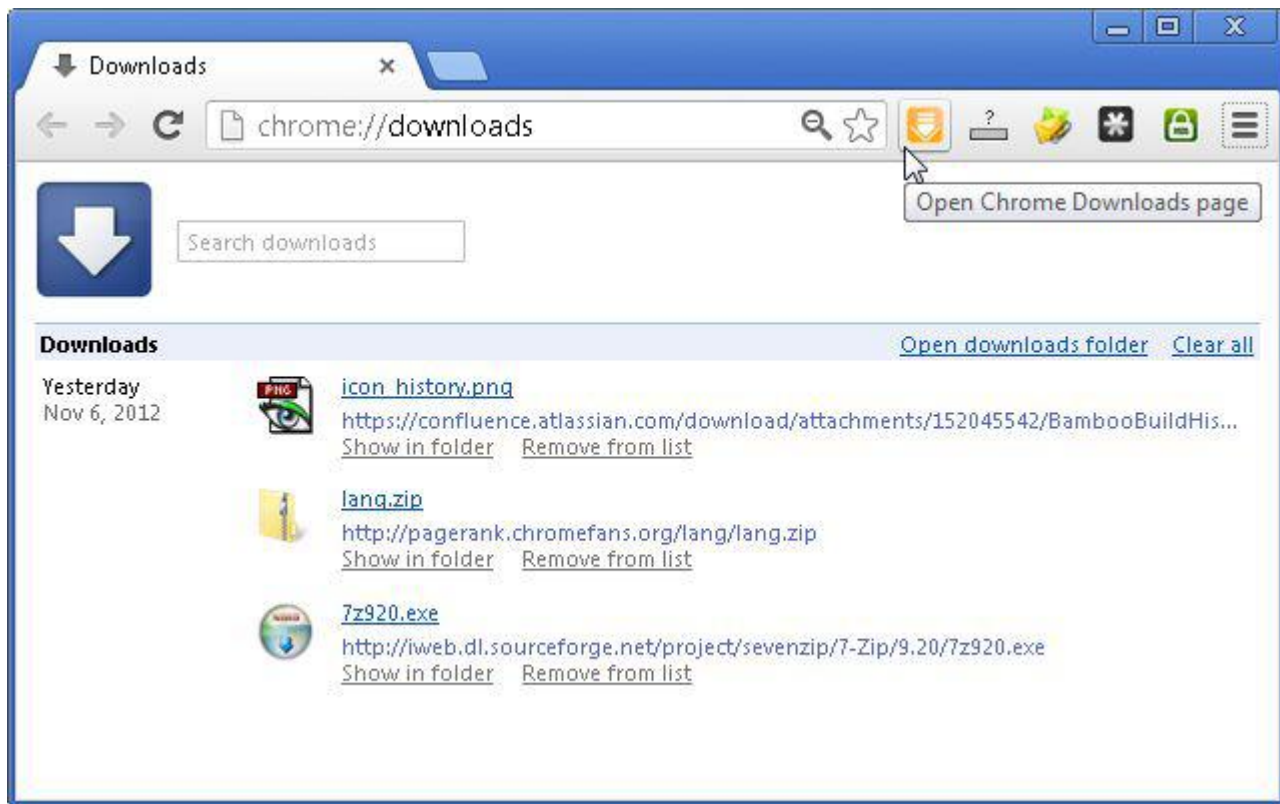
SALT-N-PEPA

LET'S TALK
ABOUT **UI** !!



Includes Bonus Track
DO YOU WANT ME
Techno Philly Mix

COMPACT DISC SINGLE




Downloads

chrome://downloads

Downloads

Search downloads

Today




100MB (1).bin

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

0 B/s - 33.8 MB of 100 MB, Paused

RESUME CANCEL




slow-auto (1).jar

<http://danbeam.org/hacks/slow-auto.jar>

This type of file may harm your computer.

DISCARD KEEP




download-1476481563.6363.txt

<http://danbeam.org/hacks/randoml.php>

Show in folder

October 5, 2016



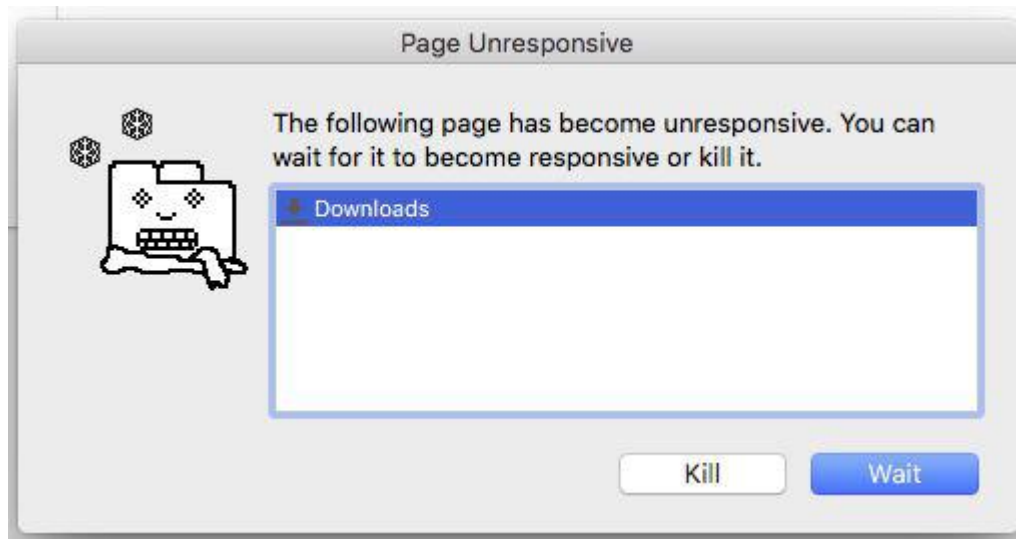
~~100MB (1).bin~~ Canceled

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

RETRY

We ported an existing code base


We didn't see performance problems, but users did



NOW WATCH THE BARS...



You're Holding it



Antenna Problems?
You're holding it wrong.



FULL BARS!



JESUSDIAZ-GIZMODO.COM

COOL HUH?

**HOLDING
IT WRONG**

**LOSES NETWORK SIGNAL
ALL THE TIME**



YOU'RE HOLDING IT

Because this does seriously happen



memeguy.com

We pretty much ignore this, saying
“We’ll make it fast later. How hard could it be?”



Old code

- Was fairly simple
- Used raw HTML/CSS/JS and had no dependencies
- Showed capped/fixed number of downloads
- Loaded/scrollled very quickly






























New code


- Was similar complexity at first glance, but...
- Had some dependencies (Polymer library + elements)
- Showed capped/fixed number of downloads
- Loaded quickly for us!

But then...

- One of the lead developers of WebKit (now Blink) filed all these bugs
 - [Polymer downloads page takes 800ms to load and then 6 seconds to show content](#)
 - [Polymer download scrollbar width code causes flood of style recalcs and layouts](#)
 - [Polymer downloads page does huge 500ms-1.8 second style recalcs and runs script for hundreds of ms when using search box](#)

They're definitely not "holding it wrong".

- ▼  Event (resize) 
- ▼  Function Call (manager.js:76) 
 -  Recalculate Style (item.js:281)
 -  **Layout (item.js:281)** 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)




Downloads

chrome://downloads

Downloads

Search downloads

Today




100MB (1).bin

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

0 B/s - 33.8 MB of 100 MB, Paused

RESUME CANCEL




slow-auto (1).jar

<http://danbeam.org/hacks/slow-auto.jar>

This type of file may harm your computer.

DISCARD KEEP




download-1476481563.6363.txt

<http://danbeam.org/hacks/randoml.php>

Show in folder


October 5, 2016



~~100MB (1).bin~~ Canceled

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

RETRY




Summary		Aggregated Details
Type	Recalculate Style	
Total Time	400.03 ms	
Self Time	400.03 ms	
Elements affected	1198	

Evaluate Script

725.10 ms

161.08 ms



1.76 s

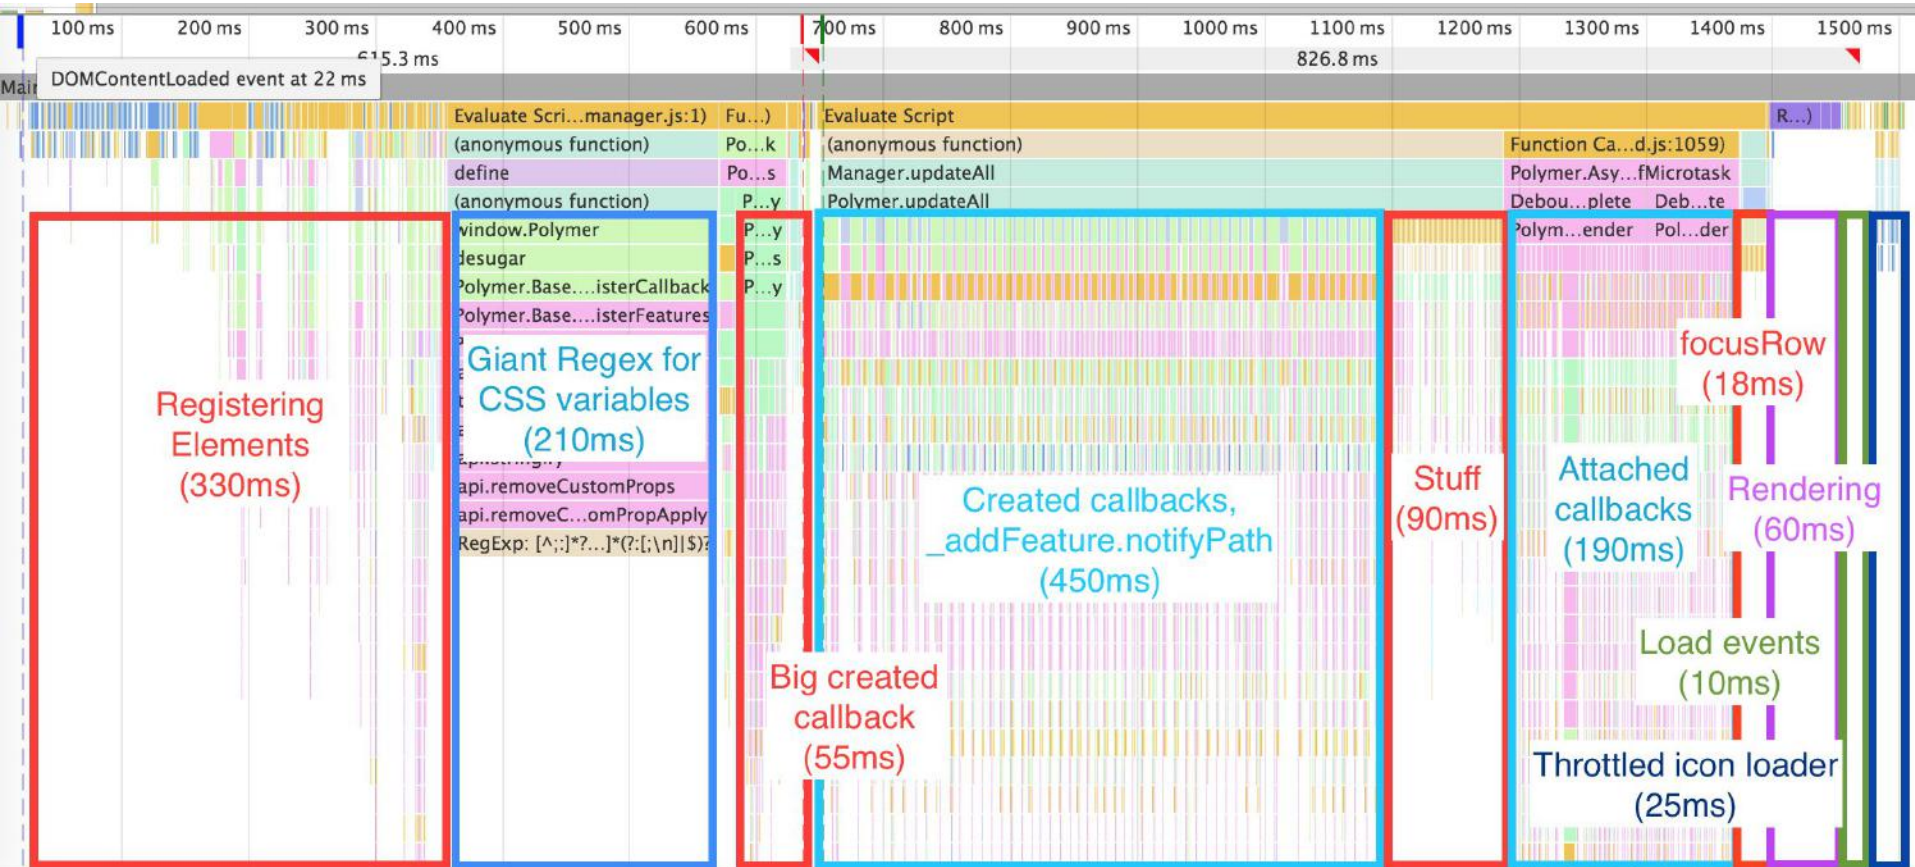
86.13 ms Scripting (Self)

1.57 ms Scripting (Children)

1.67 s Rendering

Scripting (Self)

Rendering



**So this is
definitely our
fault
at this point...**




Before you start optimizing: when is it fast enough?

Determine exit criteria

- It's fast enough when
 - † **users** can scroll at 60fps
 - It loads instantly (~130ms)
 - It works on crappy hardware

Google's web performance model: [RAIL](#)



Response

Animation

Idle

Load

How do we keep it fast?



Performance monitoring

Many ways to do this. They all boil down to:

- Agree upon/create some semi-stable testing environment[s]
- Gather inputs that might affect performance
 - poll version control for new revisions
 - use new versions of libraries
 - anything else that might affect end-user performance (i.e. simulate a bad network)
- On input changes: run the app 1 (or N) times while measuring stuff
 - CPU, GPU, power, time
- Archive and track the results over time
- Know quickly when/if your app got slow (and why)

Reproducibility is a huge problem

(which is why I was ready to ship a slow UI to users without knowing it...)

If you've ever written a UI:
congrats, you are the proud owner of a distributed system!



On the server, you **can** just throw money at the problem

You've replaced engineering ingenuity with barrels of cash. Now what?



UIs are a different kind of distributed system

They run on many **very different** machines ... that you don't have control over

- Different hardware
 - Drastically different amounts of RAM/CPU/disk
- Different host software
 - OS
 - Current load
- Different user-specific data
 - you probably won't have access to this
 - it might be hard to generate/reproduce
- Different network bandwidth/latency/reliability

Gather information from the “wild” (Remote profiling)

If you can, profile users in the most realistic way possible. Chrome has:

- `chrome://tracing`
- `inspector`
- probably lots of other stuff...

But these require us (Google) to be able to reproduce locally (on fast machines/network/with our own data) or users to send us traces.

Chrome created “UMA” to upload metrics collected on local user computers

Gather information from the “wild” (Remote profiling)

So we added remotely logged metrics for load time and other things, and guided our optimizations based on this.

High latency turnaround process (for us it was like ~weeks):

1. Make “total guess” change
2. Launch it
3. Put it front of more and more users
4. Gather data
5. Interpret data
6. Repeat

And lastly, if you can't make things actually fast...

Pretend they are.

Pretend to be fast

Trick the human brain

- Loading shims
- Progressive enhancement
- Progress bars (meh)
- Humans are only so fast
(~16.7ms frames are good enough for average eye)
- Humans detect instant at ~130ms
(ish, but this is slowly changing)



Are you actually gonna tell us how you made
Chrome's UI fast?



POLYMER
SUMMIT

JOIN THE DISCUSSION AT



UCLA

CS 130
Software Engineering

So far so good...



No! Make the **code** lazy! (oh!)

- Virtualization
 - How “optimizable” is your UI? How much is showing on the “critical path”?
- Defer resource allocation until acquisition

This takes work and generally adds complexity.



1. Reduce the

Conditional rendering

$$T = N \times K$$



POLYMER
SUMMIT

JOIN THE DISCUSSION AT



Downloads

Only 4 of these items actually fit on the screen?!
Why are we drawing 50-150?

Downloads

Today



100MB (1).bin

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

0 B/s - 33.8 MB of 100 MB, Paused

RESUME CANCEL



slow-auto (1).jar

<http://danbeam.org/hacks/slow-auto.jar>

This type of file may harm your computer.

DISCARD KEEP



download-1476481563.6363.txt

<http://danbeam.org/hacks/randomdl.php>

Show in folder

October 5, 2016



100MB (1).bin Canceled

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

RETRY

People



Dan Beam

Syncing to dbeam@google.com

Turn off



Managed by **google.com**



Sync

On - custom settings



Manage your Google Account



Chrome name and picture



Import bookmarks and settings



Autofill



Passwords



Payment methods



Addresses and more



Appearance

Themes

Open Chrome Web Store



- ☒ Open the New Tab page
- ☐ Continue where you left off
- ☐ Open a specific page or set of pages

Advanced ^

Privacy and security

Google Chrome may use web services to improve your browsing experience. You may optionally disable these services. [Learn more](#)

Allow Chrome sign-in

By turning this off, you can sign in to Google sites like Gmail without signing in to Chrome



Use a prediction service to help complete searches and URLs typed in the address bar



Use a prediction service to load pages more quickly



Use a web service to help resolve navigation errors



Safe Browsing

Protects you and your device from dangerous sites



Help improve Safe Browsing







Sends some system information and page content to Google










Automatically send usage statistics and crash reports to Google



Settings

-  People
-  Autofill
-  Appearance
-  Search engine
-  Default browser
-  On startup

Advanced

-  Privacy and security
-  Languages
-  Downloads
-  Printing
-  Accessibility
-  System
-  Reset settings

Extensions



About Chrome

Search settings

- ☒ Open the new Tab page
- ☐ Continue where you left off
- ☐ Open a specific page or set of pages

Advanced

Privacy and security

Google Chrome may use web services to improve your browsing experience. You may optionally disable these services. [Learn more](#)

Allow Chrome sign-in

By turning this off, you can sign in to Google sites like Gmail without signing in to Chrome



Use a prediction service to help complete searches and URLs typed in the address bar



Use a prediction service to load pages more quickly



Use a web service to help resolve navigation errors



Safe Browsing

Protects you and your device from dangerous sites



Help improve Safe Browsing

Sends some system information and page content to Google



Automatically send usage statistics and crash reports to Google



People



Dan Beam

Syncing to dbeam@google.com

Turn off



Managed by google.com



Sync

Open custom settings

Manage your



Chrome na



Import boo



Import bookmarks and settings

Safari



Select items to import:



Favorites/Bookmarks

Cancel

Import

Autofill



Passwords



Payment methods



Addresses and more



Appearance

Themes

Open Chrome Web Store



← Customize fonts

Font size



Minimum font size

© The quick brown fox jumps over the lazy dog



Standard font

Times ▼

16: The quick brown fox jumps over the lazy dog

Serif font

Times ▼

16: The quick brown fox jumps over the lazy dog


Sans-serif font

Helvetica ▼

16: The quick brown fox jumps over the lazy dog

Fixed-width font

About Chrome



Google Chrome is up to date
Version 72.0.3626.121 (Official Build) (64-bit)

Automatically update Chrome for all users [Learn more](#)

Get help with Chrome



Report an issue



Google Chrome
Copyright 2019 Google Inc. All rights reserved.

Google Chrome is made possible by the [Chromium](#) open source project and other [open source software](#).

Google Chrome [Terms of Service](#)



2. Reduce the cost






























$$T = N \times K$$





POLYMER
SUMMIT

JOIN THE DISCUSSION AT



- ▼  Event (resize) 
- ▼  Function Call (manager.js:76) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)
 -  Layout (item.js:281) 
 -  Recalculate Style (item.js:281)





Summary		Aggregated Details
Type	Recalculate Style	
Total Time	400.03 ms	
Self Time	400.03 ms	
Elements affected	1198	


Downloads

chrome://downloads

Downloads

Search downloads

Today




100MB (1).bin

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

0 B/s - 33.8 MB of 100 MB, Paused

RESUME CANCEL




slow-auto (1).jar

<http://danbeam.org/hacks/slow-auto.jar>

This type of file may harm your computer.

DISCARD KEEP




download-1476481563.6363.txt

<http://danbeam.org/hacks/randoml.php>

Show in folder

October 5, 2016



~~100MB (1).bin~~ Canceled

<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

RETRY


Downloads x

chrome://downloads

Downloads


Search downloads

Today




100MB (1).bin
<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>
0 B/s - 33.8 MB of 100 MB, Paused

[RESUME](#) [CANCEL](#)



slow-auto (1).jar
<http://danbeam.org/hacks/slow-auto.jar>
This type of file may harm your computer.


[DISCARD](#) [KEEP](#)



download-1476481563.6363.txt
<http://danbeam.org/hacks/randoml.php>

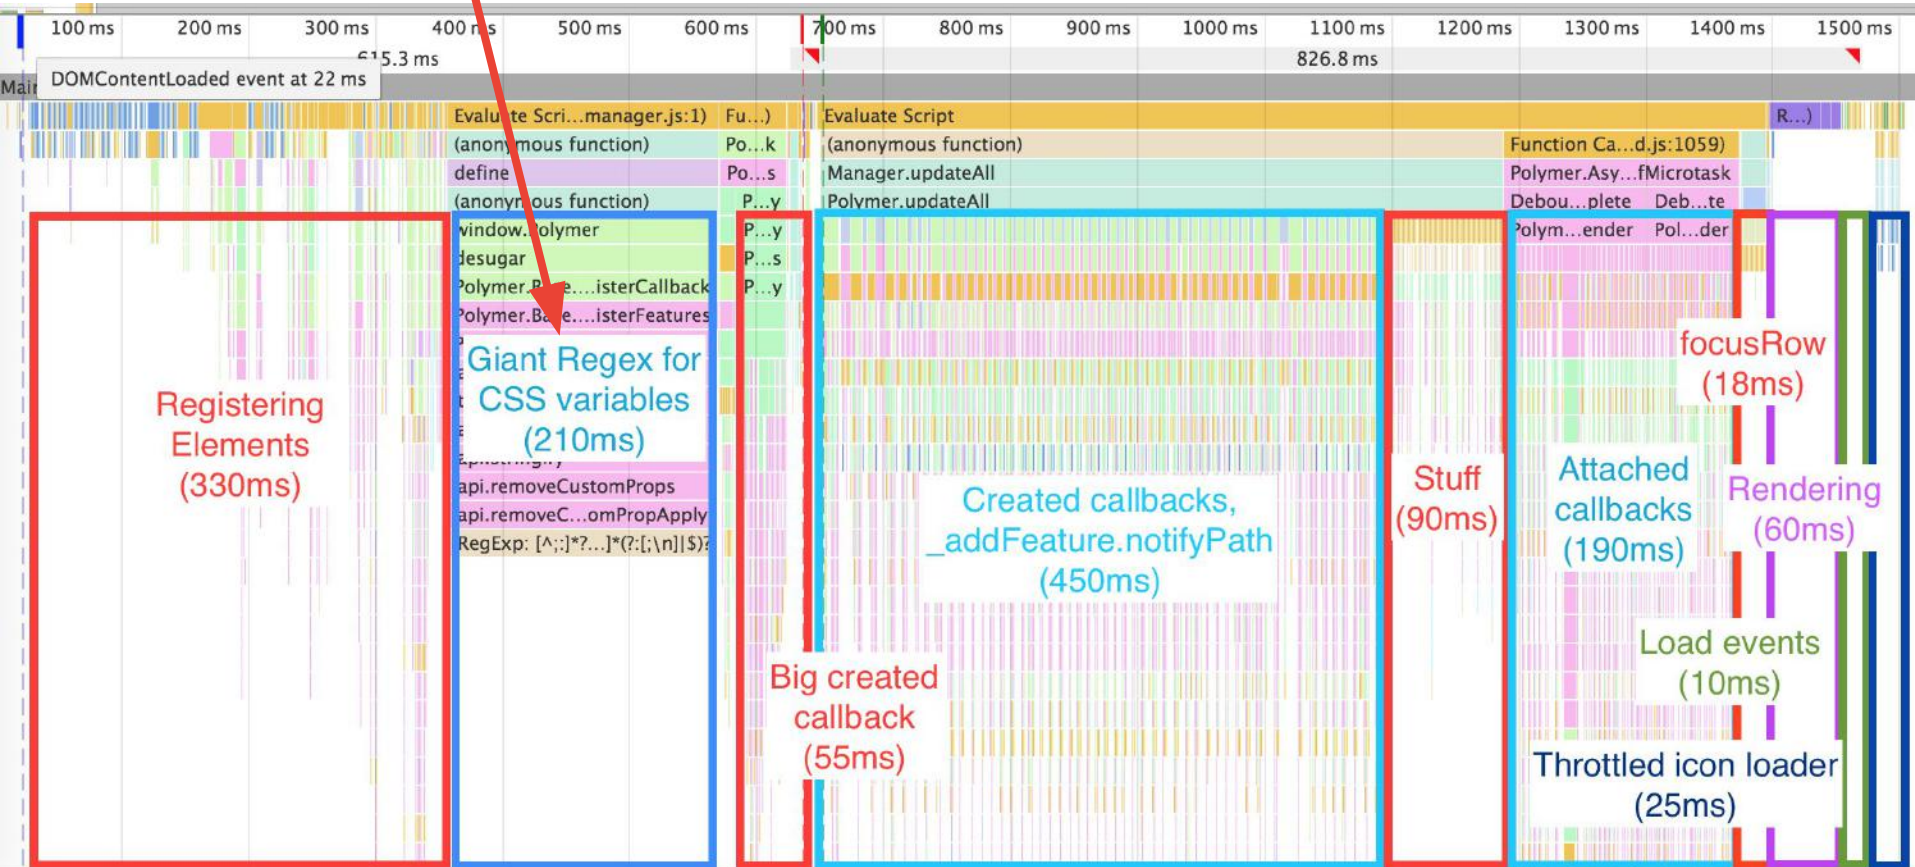
[Show in folder](#)


October 5, 2016



100MB (1).bin **Canceled**
<http://azspdcentralus.blob.core.windows.net/private/100MB.bin?sv=2015-12-11&s...>

[RETRY](#)

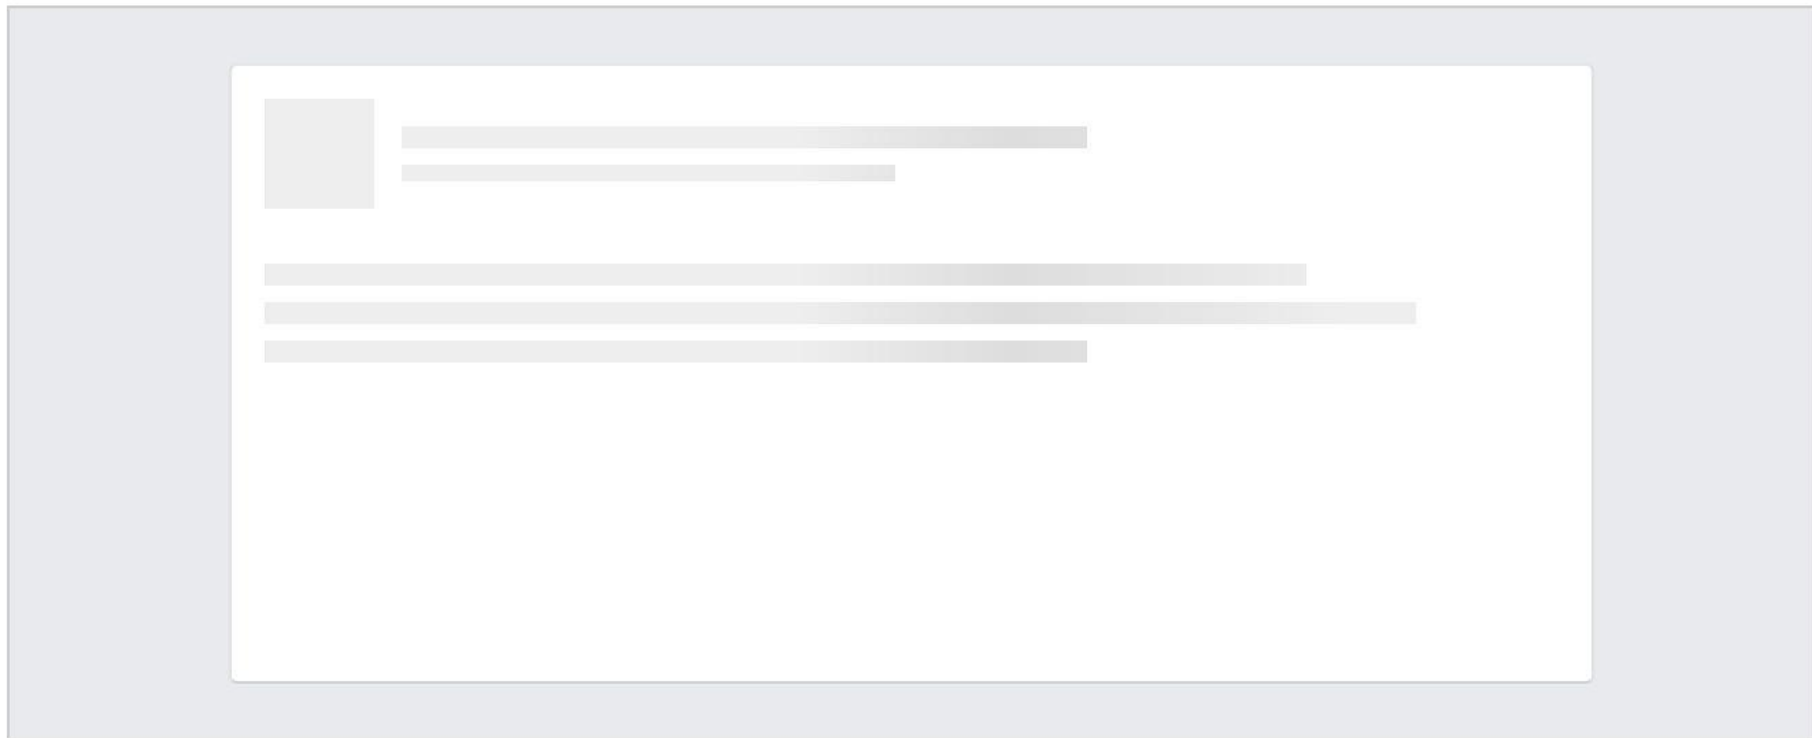






Loading...

Facebook's loading shims



Capacity Planning

Capacity Planning

You've replaced engineering ingenuity with barrels of cash. Now what?



Capacity Planning

Seems simple at first!

- In a lot of cases, your server responds linearly with traffic
 - E.g. You know that if you allocate an additional 2x memory and CPU, your server can handle 2x QPS at the same latency
- However, at the extreme ends, servers start responding non-linearly with load, which can cause all sorts of headaches

Capacity Planning -- Automation


Automation can help us out a lot.

For a lot of cases, an automated system can spin up new instances of our server as traffic increases.

Even more granular, a service like Google Cloud could increase your server's resources (e.g. memory, CPU) as traffic increases.

But how do we deal with those non-linear server responses?

Load testing case study



Common Failure Modes of Distributed Systems

A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.

-- Leslie Lamport

Common failure modes of distributed systems

Correctness bugs:

- Can often be detected in regression tests
- These systems are often data-parallel, so you can test correctness with limited resources

Performance bugs:

- Hard to catch in synthetic tests
- Typically needs full scale and representative load to properly test



Above: Common error screen around the time you were a baby

Producer-consumer rate mismatch

Even if you have admission control, a rate mismatch can be very problematic

- Best case, the back pressure propagates through your system and the system stalls or the failure is visible to the user
- Worst case, some node in the middle of your computation fails because it is overloaded




Hotspots

The load on your system can be concentrated onto a small working set.

- Results in the machines hosting that fraction of the data getting hotter than average and increasing queuing delay
- Even with load shedding, the node can “heat up” faster than you’d normally rebalance


Need to seek ways to make your



Death ray

The evil cousin of the hotspot...

- Typically happens when the load is so overwhelming that admission control can't prevent the process from failing
- However, the load doesn't cease
- When your system recovers and loads those resources onto another node, the death ray follows



Amplified rare events

Even rare things tend to be common at scale:

- Bug triggers 1 in 1M operations, running 10 ops/sec, this bug will trigger daily!

These types of failures abound. A few examples that we've observed:

- Rare checksum computation failure because of processor issue
- Remote cache restart causes tablet server failure



Debugging Distributed Systems

Debugging Distributed Systems

Once you've built a system, you often live with it for a long time.

You need tools to help you understand what the system is doing:

- At a micro scale: per job
- At a macro scale: interactions between jobs



Six Stages of Debugging

1. That can't happen.
2. That doesn't happen on my machine.
3. That shouldn't happen.
4. Why does that happen?
5. Oh, I see.
6. How did that ever work?

Status Pages

Key features:

- Each process serves over HTTP
- Gives you a local understanding of what is going on in a single process
- Can expose state of internal data or reasons for recent decisions

Especially useful status pages:

- Stacktrace of all currently live threads
- State of all currently pending RPCs

Apache Server Status for 127.0.0.1

Server Version: Apache/2.2.22 (Unix) DAV/2 PHP/5.3.3 mod_ssl/2.2.22 OpenSSL/1.0.0-fips mod_watch/4.3
Server Built: May 11 2012 16:00:00

```
Current Time: Tuesday, 21-Aug-2012 14:57:17 EDT
Restart Time: Tuesday, 21-Aug-2012 14:57:02 EDT
Parent Server Generation: 0
Server uptime: 15 seconds
Total accesses: 1 - Total Traffic: 0 kB
CPU Usage: u0 s0 cu0 cs0
.0667 requests/sec - 0 B/second - 0 B/request
1 requests currently being processed, 5 idle workers
```

[illegible]

Scoreboard Key:

"_": Waiting for Connection, "s": Starting up, "r": Reading Request, "R": Sending Reply, "k": Keepalive (read), "D": DNS Lookup, "C": Closing connection, "L": Logging, "G": Gracefully finishing, "I": Idle cleanup of worker, ".": Open slot with no current process

Srv	PID	Acc	M	CPU	SS	Req	Conn	Child	Slot	Client	VHost	Request
0-0	13856	0/0/0	W	0.0	0	0	0.0	0.00	0.00	127.0.0.1	fresh-cm.corp.interworx.com	GET /server-status HTTP/1.0
2-0	13858	0/1/1	_	0.0	6	0	0.0	0.00	0.00	127.0.0.1	fresh-cm.corp.interworx.com	GET /watch-flush HTTP/1.0

Logging

It is critical to produce useful logs.

- Logs are very useful for debugging failures
- Since you don't know when failures are going to happen, you need to log all the time

What to log:

- When and what requests arrive
- Details about erroneous conditions

What part of the code is involved

```
2013/10/30 02:21:34,177 INFO [org.jboss.as.connector.subsystems.datasources] (ServerService
2013/10/30 02:21:34,422 INFO [org.jboss.jaxr] (MSC service thread 1-2) JBAS014000: Started JA
2013/10/30 02:21:35,122 INFO [org.jboss.as.connector.subsystems.datasources] (ServerService
2013/10/30 02:21:35,525 INFO [org.jboss.as.connector.subsystems.datasources] (ServerService
2013/10/30 02:21:38,405 WARN [org.jboss.as.messaging] (MSC service thread 1-2) JBAS011600: A
2013/10/30 02:21:40,012 INFO [org.apache.coyote.http11] (MSC service thread 1-4) JBWEB003001
2013/10/30 02:21:40,207 INFO [org.apache.coyote.http11] (MSC service thread 1-4) JBWEB003000
2013/10/30 02:21:40,839 INFO [org.jboss.ws.common.management] (MSC service thread 1-1) JBWS0
2013/10/30 02:21:42,808 INFO [org.hornetq.core.server] (MSC service thread 1-4) HQ221000: liv
2013/10/30 02:21:42,811 INFO [org.hornetq.core.server] (MSC service thread 1-4) HQ221006: Wa
2013/10/30 02:21:42,717 INFO [org.jboss.as.jacorb] (MSC service thread 1-3) JBAS016330: CORBA
2013/10/30 02:21:43,511 INFO [org.infinispan.configuration.cache.EvictionConfigurationBuilder
2013/10/30 02:21:43,614 INFO [org.infinispan.configuration.cache.EvictionConfigurationBuilder
2013/10/30 02:21:44,519 INFO [org.hornetq.core.server] (MSC service thread 1-4) HQ221013: Us
2013/10/30 02:21:46,716 INFO [org.jboss.as.server.deployment.scanner] (MSC service thread 1-
2013/10/30 02:21:47,105 INFO [org.jboss.as.server.deployment] (MSC service thread 1-3) JBAS0
2013/10/30 02:21:48,104 INFO [org.hornetq.core.server] (MSC service thread 1-4) HQ221034: Wa
2013/10/30 02:21:48,104 INFO [org.hornetq.core.server] (MSC service thread 1-4) HQ221035: Liv
2013/10/30 02:21:49,045 INFO [org.jboss.as.jacorb] (MSC service thread 1-1) JBAS016328: CORBA
2013/10/30 02:21:49,829 INFO [org.jboss.as.connector.subsystems.datasources] (MSC service th
```

Local Request Tracing

Logs are useful, but they interleave events

- To remedy this, you can maintain request specific logs, often called traces
- Typically only kept for active requests, so they are less useful for retrospective debugging
- Very useful for answering the question: “What is going on right now?”

```
// Annotate a request trace.
```

```
const string& request = ...;  
if (HitCache()) {  
    TRACEPRINTF("cache hit for %s",  
                request.c_str());  
} else {  
    TRACEPRINTF("cache miss for %s",  
                request.c_str());  
}
```

Distributed Tracing

Key features:

- Integrated into each component of the whole system
- A token, representing a logical operation, is passed along each RPC
- Stats are logged by each process and a system aggregates and visualizes the data
- Tend to be verbose, so they are sampled

Types of stats:

- Latency

Resource Usage (CPU, Disk, etc)

