

1a)

- Answer:
 - For each point, we can consider 15 points on the right as a possible new closest-pair
- Proof:
 - In order for there to be a new closest pair, the newly formed pair's distance must be shorter than the current closest pair's distance
 - This current shortest distance is defined by some value d
 - In order for the new closest pair to be less than d apart from each other, both elements of the pairing must fall within d of the merge point in relation to the x-axis
 - By similar logic, the right element of the pairing must be within the left point's y-coordinate, plus or minus d (and vice versa)
 - Combining these requirements, we see we only have to analyze a $d \times 2d$ space on each side
 - Within this $d \times 2d$ space, we know that we can split this space up into $d/2 \times d/2$ boxes, such that there is only 1 point inside each of these boxes
 - Assume there were 2 points inside one of the boxes
 - The largest distance these 2 points could be from each other is $\sqrt{2d^2/4}$ or $(d * \sqrt{2})/2$
 - This contradicts the setup of our problem, as $(d * \sqrt{2})/2 < d$, yet, we defined d as the smallest distance between any points that fall within the same subproblem
 - Since the theoretical points that share a grid box are in the same subproblem, it is impossible for the distance between them to be less than d
 - This division into boxes results in 8 boxes within the significant region on either side
 - As proved above, we only need to analyze a single point within each of the boxes
 - This results in us needing to analyze 8 total points per side, or 16 total points
 - However, the point we would be comparing to must occupy one of these 16 boxes, so we only need to compare it to 15 other points

1b)

- Discussion:
 - At the beginning of the algorithm, we not only create a global list of the points sorted by x-component, we also create a global list of the points sorted by y-component
 - At each merge, we explained in part 1a why you only need to analyze points that fall within d units (by x-coordinate) of the merge point
 - As a result, we can simply iterate through our global list that is sorted by y-component to construct a local list containing only points within the current merge's slice, also sorted by y-component
 - From here, we can compare each point in the list to the closest 15 points to it and update our minimum distance accordingly
- Proof:
 - We proved in 1a that only points with an x-component within d units of the merge point can be possible candidates for closest pair
 - Using this list, we know that only 15 points are possible candidates for closest pair
 - We can simplify this by comparing each point to the next 15 points in our sorted list
 - This is because, by the time we're analyzing a point, we will have already compared it to all 15 points prior to it in the sorting since we're going in order of the sorted list
 - This covers all our bases, as we are guaranteed to have compared each point to its 15 closest neighbors by y-coordinate

2a)

- Answer:
 - Yes, we can just focus on the edges of T and the newly added weights
- Proof:
 - MST theorem states that, when partitioning the graph into 2 partitions, the minimum weight edge that crosses these partitions is part of the MST
 - Take the current tree T as one partition and the vertex X alone in its own partition
 - By MST theorem, the minimum weight edge that connects T to X must be a part of the MST
 - As T is already an MST, adding this newly processed edge to T guarantees a new MST containing X by MST theorem

2b)

- Answer:
 - T is a minimum spanning tree of G
- Proof:
 - Let the number of nodes in the graphs G and G' be represented by n
 - This means that there must be $n - 1$ edges in the MST for either graph, as all nodes are part of the tree and, by definition, the tree cannot have any cycles
 - Due to this, if we decrease the weight of each edge by some constant k, we're necessarily decreasing the total weight of the MST of that graph by $k(n - 1)$, as all $n-1$ edges in the MST have had their weight reduced by k
 - This tells us that if the MST of T' is some weight W, the MST of T must be $W - k(n - 1)$
 - Since the difference between G' and G is that every edge in G has k less weight than its corresponding edge in G', an MST made up of edges in G will be of weight $k(n - 1)$ less than the corresponding MST in G'
 - This follows exactly with our statement above, as the total weight of the MST in G will be the total weight of the MST in G' minus $k(n - 1)$, or $W - k(n - 1)$
 - By sanity check, this proof makes sense as each edge is simply decremented by constant, so the MST in either graph shouldn't change

3)

- Algorithm:
 - Initialize a variable L to 0
 - Initialize a variable R to the size of B - 1
 - While L is less than R:
 - Calculate an index mid using $\text{floor}([L + R] / 2)$
 - If the value at index mid is greater or equal to the value at index L:
 - Set L to mid + 1
 - Else:
 - Set R to mid
 - Save the value of L in a variable M
 - Reset L to 0 and R to the size of B - 1
 - While L is less than or equal to R:
 - Calculate an index temp using $\text{floor}([L + R] / 2)$
 - Set a new variable mid equal to (M + temp) modulo (the size of B)
 - If the value at index mid is equal to X:
 - Return true
 - Else if the value at index mid is less than X:
 - Set L equal to temp + 1
 - Else:
 - Set R equal to temp - 1
 - Return false
- Proof:
 - The overall structure of the algorithm is to find how much the array was shifted by, then using that value to modify a basic binary search of the array
 - Since A is sorted, the amount that B was shifted by is equivalent to the index of the minimum element in B
 - This minimum element is differentiated, as it is the only element where the element that directly precedes it is greater than itself
 - Base Case:
 - The size of B is 1
 - If this element is X, then X is in B, otherwise it's not
 - The initial shift calculation should tell us that the array was not shifted ($M = 0$)
 - Based on initialization, the loop containing the first binary search will not run
 - L's initial value of 0 will be saved to M
 - Shift calculation passes the base case
 - The following binary search will begin
 - This search should simply check if the element in B is in fact X
 - The loop will run 1 time
 - A mid value will be calculated, and will remain unadjusted, as there have been no shifts

- The first value will be checked, and true will be returned if the value matches X
 - If not, the loop will terminate and return false
 - This is the expected behavior
 - Base case passed
- Inductive Step (Shifting):
 - Assume: The algorithm has executed so that there are n elements to search remaining
 - Prove: The current step will return a subproblem that contains the pivot point
 - A midpoint index will be calculated by a simple average
 - The algorithm will do one of 2 things:
 - It will set L to mid + 1:
 - For this to happen, the algorithm must have found that the value at mid is \geq to the value at L
 - This tells us that we're still in an increasing part of the array
 - Since the array is non-decreasing from L to mid, we know that the minimum value (pivot point) cannot be located in that subarray
 - Therefore, setting L to the index beyond that subarray is safe, as we know that the minimum must be in the right subarray
 - This is the correct behavior
 - It will set R to mid:
 - For this to happen, the algorithm must have found that the value at mid is $<$ the value at L
 - Based on this, we know that, at some point in the subarray between L and mid, the values in the array decrease
 - Based on our reasoning from the beginning of the proof, we know that this tells us our minimum must be in this subarray
 - As a result, we move the right bound of R to the end of the subarray
 - This is the correct behavior
 - Both situations exhibit the correct behavior
 - We can inductively extend this logic to say that our resulting subarray contains the minimum element, which means we know how much the array was shifted by
 - Inductive step complete
- Inductive Step (Searching):
 - Assume: The loop has executed so that there are n elements to search through

- Prove: The current step will return a subproblem that contains the pivot point
 - The algorithm will calculate a temporary midpoint index
 - This temporary midpoint index will be adjusted to reflect the index of the midpoint if the array were sorted
 - This is done by adding the amount that B was shifted (M) to the temporary midpoint index and then taking the remainder of that calculation to re-index it into the array B
 - By doing this, each temporary midpoint index we can calculate will correspond to a distinct and valid final midpoint index
 - From here, the algorithm will do one of 3 things:
 - It will find that the value at mid is equal to X
 - At this point, we know X is in fact in B
 - True is returned
 - This is the correct behavior
 - It will find that the value at mid is less than X:
 - Since the array is sorted, we then know that, if X exists, it is located in the right subproblem
 - This is adjusted for by moving the left boundary to the right of mid
 - This is the correct behavior
 - It will find that the value at mid is greater than X:
 - Since the array is sorted, we then know that, if X exists, it is located in the left subproblem
 - This is adjusted for by moving the right boundary to the left of mid
 - This is the correct behavior
 - All 3 situations result in the expected behavior
 - This tells us that, upon completion of the current iteration, the result will have been found, or a subproblem containing the possible solution will have been found
 - We can inductively extend this logic to further subproblems until the target X has been found or there are no more elements to search
 - We will know the answer by the end of execution
- Inductive step complete
 - Proof by induction complete
- Time Complexity: $O(\log n)$
- Time Complexity Proof:
 - Outside of the 2 while loops, only constant time operations exist
 - Both while loops consist of binary search, where we can use the equation $T(n) = T(n / 2) + C$

- Both binary searches are made up of a constant number of calculations and comparisons, resulting in the C term
- Both binary searches move to the next iteration with a subproblem that is approximately half the size of the previous subproblem, resulting in the $T(n / 2)$ term
- If we expand the equation, we see that it is equal to $T(n) = T(n / 2^i) + Ci$
- We know that searching an array of size 1 is $O(1)$, so we can say that $n / 2^i = 1$
 - This means $i = \log n$
- Plugging back in, we have $T(n) = 1 + C \log n$
- This reduces to tell us our time complexity is $O(\log n)$

4)

- Algorithm: // Assume n is the size of A
 - If n is less than 4:
 - Return some failure condition
 - Initialize 4 arrays, a1, a2, a3, and a4, each of size n + 1, with some value representing negative infinity
 - Initialize an iterator i to n - 1
 - While i is greater than or equal to 0: // **Get max vals of A[s]**
 - Set a1[i] equal to the maximum of a1[i + 1] and A[i]
 - Decrement i by 1
 - Set i to n - 2
 - While i is greater than or equal to 0: // **Get max vals of A[s] - A[r]**
 - Set a2[i] equal to the maximum of a2[i + 1] and a1[i + 1] - A[i]
 - Decrement i by 1
 - Set i to n - 3
 - While i is greater than or equal to 0: // **Get max vals of A[s] - A[r] + A[q]**
 - Set a3[i] equal to the maximum of a3[i + 1] and a2[i + 1] + A[i]
 - Decrement i by 1
 - Set i to n - 2
 - While i is greater than or equal to 0: // **Get max vals of A[s] - A[r] + A[q] - A[p]**
 - Set a4[i] equal to the maximum of a4[i + 1] and a3[i + 1] - A[i]
 - Decrement i by 1
 - Return a4[0]
- Proof:
 - a1 holds the max values of A[s] for each index, a2 holds the max values of A[s] - A[r] for each index, a3 holds the max values of A[s] - A[r] + A[q] for each index and a4 holds the max values of A[s] - A[r] + A[q] - A[p] for each index
 - Proof by induction:
 - For each step in the first loop, 2 things can happen
 - The entry into a1 is the current value of A:
 - This means that the current value of A is the maximum value possible thus far
 - This is confirmed by comparison
 - Correct behavior
 - The entry into a1 is not the current value of A
 - This means that the current value of A is not the maximum value possible thus far
 - This tells us the previous maximum is the optimal solution for this entry
 - This is confirmed by comparison
 - Correct behavior
 - For each step in the second loop, 2 things can happen
 - The entry into a2 is the optimal value of A[s] minus the current value of A:

- This means that the current value of $A[s] - A[r]$ is the maximum value possible thus far
 - This is confirmed by comparison
 - Correct behavior
- The entry into a2 does not involve the current value of A
 - This means that the current value of $A[s] - A[r]$ is not the maximum value possible thus far
 - This tells us the previous maximum is the optimal solution for this entry
 - This is confirmed by comparison
 - Correct behavior
- For each step in the third loop, 2 things can happen
 - The entry into a3 is the optimal value of $A[s] - A[r]$ plus the current value of A:
 - This means that the current value of $A[s] - A[r] + A[q]$ is the maximum value possible thus far
 - This is confirmed by comparison
 - Correct behavior
 - The entry into a3 does not involve the current value of A
 - This means that the current value of $A[s] - A[r] + A[q]$ is not the maximum value possible thus far
 - This tells us the previous maximum is the optimal solution for this entry
 - This is confirmed by comparison
 - Correct behavior
- For each step in the fourth loop, 2 things can happen
 - The entry into a4 is the optimal value of $A[s] - A[r] + A[q]$ minus the current value of A:
 - This means that the current value of $A[s] - A[r] + A[q] - A[p]$ is the maximum value possible thus far
 - This is confirmed by comparison
 - Correct behavior
 - The entry into a2 does not involve the current value of A
 - This means that the current value of $A[s] - A[r] + A[q] - A[p]$ is not the maximum value possible thus far
 - This tells us the previous maximum is the optimal solution for this entry
 - This is confirmed by comparison
 - Correct behavior
- Each step in the algorithm exhibits the correct behavior
- The value in $a4[0]$ represents the maximum value of the target expression, given access to the entire array, by definition of our problem construction
- Proof by induction complete

- Time Complexity: $O(n)$
- Time Complexity Proof:
 - The initialization of the the 4 arrays takes $O(n)$ time
 - Each of the 4 loops iterates through each element of the array, performing constant time comparisons on each
 - Since these loops are organized sequentially, they all contribute to an $O(n)$ time
 - The overall runtime of the algorithm is $O(n)$

5)

- Proof:
 - Hamiltonian path is NP-complete (Y)
 - Prove that ST-Hamiltonian path is also NP-complete (X)
 - To do this, we must prove that Hamiltonian path is polynomial-time reducible to ST-Hamiltonian path
 - Assume we have algorithms that solve both problems
 - Hamiltonian path takes in an input of a graph
 - Take this graph, and extract all ${}_nC_2$ combinations of start point and end point
 - ${}_nC_2$ is on the order of n^2 , so this is a polynomial time transformation
 - This is every possible combination of S and T in the graph
 - We can now proceed to pass each graph through our algorithm for ST-Hamiltonian path
 - If any of these inputs are found to have an ST Hamiltonian path, the graph also has a general Hamiltonian path
 - This means we can use a logical statement with order n^2 comparisons to transform the output of ST Hamiltonian path to Hamiltonian path
 - We have shown that Hamiltonian path is polynomial-time reducible to ST-Hamiltonian path
 - Since Hamiltonian path is polynomial-time reducible to ST-Hamiltonian path, it follows that ST Hamiltonian path must be NP-complete, otherwise it would be possible to use a series of polynomial-time transformations to transform the NP-complete Hamiltonian path problem into ST Hamiltonian path, which would result in Hamiltonian path being solvable in polynomial time

6)

- **Algorithm: // Assumes that such an assignment is possible**
 - Create a source node S and a sink node T
 - Create N nodes to represent each team
 - For each node N_i :
 - Duplicate the node into N_1 and N_2 and connect them from N_1 to N_2 with an edge with capacity t_i to reflect the size of team i
 - Create an edge with infinite capacity from S to each of the N_1 nodes
 - Create M nodes to represent each table
 - For each node M_j :
 - Duplicate the node into M_1 and M_2 and connect them from M_1 to M_2 with an edge with capacity c_j to reflect the number of chairs at table j
 - Create an edge from each node N_{i2} to each node M_{j1} with a capacity of 1
 - Create an edge with infinite capacity from each of the M_2 nodes to T
 - Run Ford-Fulkerson on the network
- **Proof:**
 - Our goal is to represent a matching of players to tables in such a way that no 2 players from the same team are at the same table
 - We enforce the number of players on each team by duplicating the nodes that represent teams, and creating an edge between them with a capacity equal to the number of players on that team
 - This results in us being able to artificially create a bottleneck on each team's node based on the number of players that team has
 - We enforce that each player can only be assigned to a table that none of their other teammates have been assigned to, by creating an edge of capacity 1 between each team and each table
 - In a high-level view of the problem, this essentially means each team can only send 1 player to a table
 - We enforce the number of chairs at each table by duplicating the nodes that represent tables, and creating an edge between them with a capacity equal to the number of chairs at that table
 - We then run Ford-Fulkerson, which will terminate when there are no augmenting edges in the residual network
 - Since there are no augmenting edges in the residual network, it must follow that there exists a cut such that the flow is equal to the capacity of the cut
 - This must be the case because we can create a cut where all nodes reachable from S are in the S partition and all other nodes are in the T partition
 - Since there are no augmenting edges, S and T are disconnected, assuming saturated edges are removed
 - By conservation of flow, all flow that leaves S must end up at T
 - This means that the capacity of the cut is made up of all saturated edges
 - Combining this with our statement on conservation of flow, the capacity of the cut must be equal to the flow in the network

- Since there exists a cut such that the flow is equal to the capacity of the cut, this flow must be the max flow
 - This is because, given a capacity of a cut, the flow in the network must be less than or equal to that capacity
 - Therefore, since we found a flow that is equal to the capacity of a cut, we know that no larger flows can exist, otherwise they would be larger than the capacity of the cut
- Since this flow is the max flow, we know that our algorithm has finished execution correctly, as players will all be assigned
- Time Complexity: $O(\sum(t_i) * (N + M) + NM)$
- Time Complexity Proof:
 - Creating all of the nodes requires creating $2N + 2M + 2$ nodes, which results in an $O(N + M)$ runtime
 - To connect source to team nodes and table nodes to sink requires $O(N + M)$ runtime
 - For all N team nodes, we must create an edge to each of the M table nodes, resulting in an $O(NM)$ runtime
 - Running Ford-Fulkerson requires $O(f(N + M))$ runtime
 - In this case, f is bounded by the minimum of the summations of t_i and c_j
 - Since we assume this assignment is possible, we know there are at least as many players as chairs, so $t_i \leq c_j$
 - As a result, Ford-Fulkerson requires $O(\sum(t_i) * (N + M))$ runtime
 - Therefore, the overall time complexity of the algorithm is $O(\sum(t_i) * (N + M) + NM)$