

1)

- Algorithm:
  - Construct a BFS tree starting from node  $s$  as the root
  - Set up an array of size  $n + 1$ , initializing each element to 0
  - Iterate through the BFS tree, calculating the level of each node by keeping an iterator to count how many edges have passed by
    - For each node, use this level as an index into the array and increment the accessed value by 1 (this keeps track of how many nodes there are per level)
  - Iterate through the array, checking for a value that is exactly 1
  - The index of this value,  $i$ , represents the level that the single node is on, so iterate through the BFS tree once more on an arbitrary path, passing exactly  $i$  edges
  - Return the node you land on
- Proof:
  - Construct a BFS tree from  $s$  to  $t$
  - The first claim is that there is a level in this BFS tree that contains only 1 node
    - By the problem statement, we know there will be at least  $n / 2$  levels between the root  $s$  and the level containing  $t$
    - Let's say there are at least 2 nodes per level
    - If this were true, the levels between  $s$  and  $t$  would contain a total of  $2(n / 2) = n$  nodes
    - This is impossible since  $s$  and  $t$  are not found in the levels between  $s$  and  $t$ , therefore there would have to be  $n + 2$  nodes in the graph
    - By the problem statement, we know there are only  $n$  nodes in the graph, so we have reached a contradiction, there must be at least 1 level containing only 1 node
  - The second claim is that the existence of a level with only 1 node implies that a node  $v$  exists that can break any  $s$ - $t$  path
    - We now know there exists some level that only contains node  $v$ , and that this level lies between the levels containing  $s$  and  $t$
    - Let's say this node  $v$  is at some level  $n$
    - Deleting node  $v$  means we cannot access any level from  $n$  onwards
    - By the definition of a BFS tree, the path to  $t$  must contain an edge leading through level  $n$ , otherwise  $t$  would have been in an earlier level of the BFS tree
    - This means that the path from  $s$  to  $t$  must contain an edge from level  $n - 1$  to  $n$

- However, the only node in  $n$  is  $v$ , which is now deleted, so any possible path from  $s$  to  $t$  is now broken
- Time Complexity:  $O(n + m)$
- Time Complexity Proof:
  - Construction of a BFS tree requires visiting each node and edge in the graph once and performing a constant time operation on each, resulting in an  $O(n + m)$  runtime
  - Allocating and initializing the array requires  $O(n)$  time
  - Iterating through the BFS tree and calculating each level also requires an  $O(n + m)$  runtime for the same reasons listed above
  - Iterating through the array has a worst case runtime of  $O(n)$ , as that is the maximum number of levels we will have to check
  - Iterating through the BFS for the third time once again has a runtime of  $O(n + m)$
  - Overall, this algorithm has a runtime of  $O(5n + 3m)$ , which can simply be reduced to  $O(n + m)$

2)\*

- Proof:
  - Assume that  $G$  is not the same graph as  $T$
  - In order for this to be true, there must be some edge  $e = (i, j)$  in  $G$  that is not in  $T$
  - There are 2 possibilities for this edge:
    - $e$  would be located within a level of  $T$ 
      - Let  $T$  be the result of the BFS and  $T'$  be the result of the DFS
      - Assuming this is the case, this means that  $i$  is an ancestor of  $j$
      - By the definition of a DFS, if  $i$  was indeed an ancestor of  $j$ ,  $j$  would be located on a lower level than  $i$ , since all ancestors in a DFS tree are guaranteed to be on a higher level than their child nodes
      - This implies  $T'$  is not equal to  $T$ , which is a contradiction of the problem statement
      - As a result, such an edge  $e$  could not be present in the graph  $G$
    - $e$  would skip over a level of  $T$ 
      - Let  $T$  be the result of the DFS and  $T'$  be the result of the BFS
      - Assuming this is the case, this means that  $i$  and  $j$  differ by more than exactly 1 level in  $T$
      - By the definition of BFS, this means that, in  $T'$ ,  $i$  and  $j$  are exactly 1 level apart or on the same level
        - This is because in a BFS, if there is an edge between 2 nodes, either one will discover the other or they will both be discovered in the same level
      - This implies that  $T$  is not equal to  $T'$ , which is a contradiction of the problem statement
      - As a result, such an edge  $e$  could not be present in the graph  $G$
  - Since both possibilities of edge  $e$  cannot exist, we have proven by contradiction that  $G = T$  if  $T$  is the BFS tree and DFS tree of  $G$

3)

- Algorithm:
  - Iterate through each of the triples, creating an undirected edge between each computer pair  $C_i$  and  $C_j$ , assigning the edge a weight of  $t_k$  to create a graph  $G$
  - Starting from the root  $C_i$  with time  $t = 0$
  - While there are still nodes to search in  $G$ :
    - For the children of the current node:
      - If the weight of the edge leading to the child node is greater than or equal to the node's time value:
        - If the child node is explored:
          - If the new edge's weight is less than the child node's time pairing:
            - Replace the time pairing with the edge weight
          - Else:
            - Skip the node
        - Else:
          - Push the node into the list of nodes to search, paired with the weight of the edge as its time value
          - Create an edge between the current node and the child node
          - Mark the node as explored
      - Else:
        - Skip the node
  - For all nodes in the BFS tree:
    - If the current node is  $C_j$ :
      - Return true
  - Return false
- Proof:
  - The construction of the graph is straightforward
    - We define our graph as a collection of computers (nodes) that are connected by an edge if they communicate at some time  $t$
    - This graph can be undirected, since the virus can be passed as long as the computers communicate, it doesn't matter which communicates with which
    - The time  $t$  at which they connected is stored for later use
  - Now, we must prove that the algorithm's next step finds all computers that  $C_i$  communicates with, both directly and indirectly
    - Assume there is a computer  $C_x$  that communicates with  $C_j$ , but is ignored by the algorithm
    - There are 2 cases:
      - $C_x$  was in direct communication with  $C_j$ :
        - For this to be true,  $C_i$  must be the parent of  $C_x$
        - Since the algorithm initializes  $C_i$ 's time as 0, all other time

- values are greater than or equal to  $C_i$ 's
  - This means that an edge leading from  $C_i$  to  $C_x$  would automatically be valid, and  $C_x$  would be analyzed by the algorithm
  - The case has been contradicted
- $C_x$  was in indirect communication with  $C_i$ :
  - The direct communication case was the base case
  - Assume the graph is correctly constructed up until  $C_x$  is analyzed
  - By definition, the edge leading to  $C_x$  must have a greater weight than its parent node's time value, otherwise  $C_x$  would not have received the virus based on our problem setup
  - If this were true, the algorithm would push  $C_x$  onto the list of nodes to be analyzed
  - This case has been proven incorrect by induction
- Both cases have been contradicted, meaning it is impossible for there to be a computer  $C_x$  that should've been accessed but wasn't
- Time Complexity:  $O(m + n)$
- Time Complexity Proof:
  - The creation of the graph  $G$  involves performing multiple constant time operations on each triple, which reduces down to an  $O(m)$  time since there are  $m$  triples
  - The construction of a BFS tree requires visiting each node and edge in the worst-case, resulting in an  $O(m + n)$  runtime
  - The search of the BFS tree requires a worst-case time of  $O(m + n)$  to analyze each node
  - The overall runtime can then be reduced to  $O(m + n)$

4b)\*

- Counterexample:
  - This statement is false
  - Assume a simple case:
    - All weights are positive, the graph is undirected, and the graph is connected
    - The path  $P$  leads from  $s$  to  $t$  and has weight 4
    - An alternate path exists that goes through some node  $x$  with edge weights of 3 and 3
    - Assuming these are the only nodes and edges in the graph, by definition,  $P$  is a minimum-cost path, since  $P$  has a weight of 4 and the path through  $x$  is of total weight 5
    - However, when squared the path  $P$  takes on a weight of 16, while the path through  $x$  takes on a weight of 13
    - $P$  is clearly no longer the minimum length path, as  $16 > 13$
    - The statement must be false since it clearly doesn't apply to all possible cases

5)\*

- Algorithm:
  - Initialize a counter variable to 0 and an element variable
  - For all elements in the array:
    - If the current counter is 0:
      - Set the element variable to the current element
      - Set counter to 1
    - Else if the current element is equal to the element variable:
      - Increment the counter variable by 1
    - Else:
      - Decrement the counter variable by 1
  - Set the counter variable to 0
  - For all elements in the array:
    - If the current element is equal to the element variable:
      - Increment counter by 1
  - Return if the counter is strictly greater than  $n / 2$
- Proof:
  - The second for loop is self-explanatory
    - We're simply checking to see if the result of the first for loop is a majority through brute force iteration
  - Therefore, we must simply prove that the result returned from the first for loop is the best possible option for a majority element
  - There are 3 cases that occur during the algorithm's runtime:
    - There's no majority vote:
      - There is no "best" element to be returned, since all of them are equally not the majority
      - The absence of a majority is caught by the second for loop, as the element returned from the first for loop is not a majority
      - The correct result is returned
    - The counter never returned to 0
      - In order for the counter to not return to 0, at any given point in the array traversal, it must have incremented more times than it has decremented
      - For this to be true, the first element encountered in the algorithm's runtime must be the majority element
      - Since the counter never returns to 0, the element variable never changes from the first element
      - The first element is therefore returned from the first for loop
      - The correct result is returned
    - The counter returned to 0 at some point
      - Each time the counter returns to 0, the algorithm essentially resets
      - This means that when the counter returns to 0, the majority element of the whole array must still be the majority element of the elements in the remaining subarray

- If we can prove this to be true, we've inductively proved that the algorithm will return the correct value for this case
- Assume there are  $x$  occurrences of the majority element in the array
- By the definition of a majority element  $x > n / 2$
- Imagine we have searched  $i$  elements when the counter reaches 0
- Assuming a lower bound case for the number of majority elements remaining, there are  $x - i$  majority elements left in the subarray and  $n / 2 - i$  total elements left in the subarray
- Since  $x > n / 2$ , it is implied that  $x - i > n / 2 - i$ , which means the majority element is still the majority element in the subarray
- In the base case where there is a single element left in the array (the majority element by definition), the algorithm increments counter to 1 and returns the correct value, so the base case is satisfied
- The correct value is guaranteed to be returned
- Time Complexity:  $O(v)$
- Time Complexity Proof:
  - The first pass iterates through all elements of the array once, performing constant time operations on each of them, resulting in an  $O(v)$  complexity for the first for loop
  - The second pass iterates through all elements of the array once, performing a check and possible iteration, both of which are constant time, resulting in an  $O(v)$  complexity for the second for loop
  - This reduces down to an  $O(v)$  runtime for the entire algorithm



6a)\*

- Algorithm:
  - If there are no edges:
    - Return 0
  - Arbitrarily order all edges in a list  $x$
  - Initialize a size variable to  $e$
  - While the size variable is greater than 0:
    - For all elements in  $x$ :
      - While new permutations beginning with  $x$  of size  $e$  still exist:
        - Recursively find an unused permutation of  $x$  beginning with the current element
        - If the permutation of  $x$  represents a valid path in the DG
          - Return the current size variable
        - Else:
          - Mark the current permutation as used
      - Decrement the size variable
    - Return 0
- Proof (informal since the previous problem explicitly asked for a proof and this problem doesn't ask for it):
  - The interior of the algorithm finds each possible permutation of the edges presented in the problem
  - The algorithm therefore checks every possible path starting from longest possible ( $e$ ) to shortest (0)
  - Therefore, the first permutation that represents a valid path must be the longest length path in the entire graph
- Time Complexity:  $O(e!)$
- Time Complexity Proof:
  - The algorithm generates every permutation of  $e$  edges, and checks if the path generated by each permutation is valid
  - The check of each path takes at most  $O(e)$  time
  - The check of every permutation takes  $O(e!)$  time
  - If the first check fails, the algorithm goes to check  $(e - 1)!$  elements by the definition of an  $r$ -permutation
    - This continues from  $n, n - 1, n - 2, \dots, 1$
  - The  $e!$  term dominates the time complexity, so the algorithm is overall  $O(e!)$

6b)

- Algorithm:
  - Perform a topological ordering of the given graph
  - Starting from the beginning of the topological ordering:
    - If the node has an indegree of 0:
      - Pair the node and out-edges with the value 0
    - Else:
      - Pair the node with the maximum value of its in-edges + 1

- Initialize a max value variable to 0
- For each node in the topological ordering:
  - If the node's paired value is greater than the max value:
    - Update the max value to the current node's value
- Return the max value variable
- Proof (informal since the previous problem explicitly asked for a proof and this problem doesn't ask for it):
  - The guarantee that this graph is acyclic tells us that the longest path must begin at a source and end at a sink, as any other start/end would have an indegree/outdegree that could be used to extend the path
  - The use of topological sort allows us to generate all possible paths from source(s) to sink(s)
  - By virtue of topological sort's design, it is then possible to iterate through the sorted list of nodes and track local maximum path lengths
    - This is because in a topological sort, we don't have to worry about anything prior to the current node, we know for a fact everything until the current point is updated
  - By tracking lengths, we can then reach the sink(s) and save maximum values obtained from them, leading to the right output
- Time Complexity:  $O(n + e)$
- Time Complexity Proof:
  - Topological sort takes  $O(n + e)$  time
  - The first loop iterates through the graph, accessing each node once and most edges twice, resulting in an  $O(n + e)$  runtime
  - The final loop runs through each node, resulting in an  $O(n)$  runtime
  - The overall runtime of the algorithm can be reduced to  $O(n + e)$